

Dynamic Prototype Routing for Lightweight Image Super-Resolution

Anonymous CVPR submission

Paper ID ****

Abstract

001 *Lightweight single image super-resolution (SR) seeks a fa-*
002 *vorable trade-off between reconstruction quality and model*
003 *cost for resource-constrained platforms. Content-aware*
004 *routing, which groups tokens by semantics rather than spa-*
005 *tial position, has recently advanced this trade-off, but its*
006 *prevailing realization—the Token Aggregation Block (TAB)*
007 *of CATANet—relies on cross-batch exponential-moving-*
008 *average routing centers, discards the confidence of each*
009 *token-to-center assignment, and entangles prototype gen-*
010 *eration with confirmation. We propose Dynamic Prototype*
011 *Routing (DPR), a drop-in replacement for TAB that keeps*
012 *the backbone, the tensor interface, and the lightweight bud-*
013 *get unchanged. DPR generates semantic prototypes directly*
014 *from the current tokens through a soft assignment, confirms*
015 *them with a learnable query bank decoupled from genera-*
016 *tion, matches tokens back with a learnable temperature, and*
017 *propagates an explicit per-token confidence through order-*
018 *ing, gating, and a soft fallback; a usage-balancing regu-*
019 *larizer keeps prototype use spread out. The resulting net-*
020 *work, DPRNet, is trained on DIV2K and evaluated on five*
021 *standard benchmarks at $\times 2 / \times 3 / \times 4$. DPRNet matches*
022 *or slightly surpasses the strongest content-routing baseline*
023 *across benchmarks at a comparable parameter and com-*
024 *putation budget (e.g. 26.89 dB on Urban100 and 45.77 G*
025 *Multi-Adds at $\times 4$, the latter below CATANet), while expos-*
026 *ing an interpretable, confidence-aware routing mechanism*
027 *whose prototype usage is measurably more balanced.*

028 1. Introduction

029 Single image super-resolution (SR) aims to reconstruct a
030 high-resolution (HR) image from its low-resolution (LR)
031 observation, and it underpins a wide range of applica-
032 tions from mobile photography to display rendering. Since
033 the seminal deep model SRCNN [3], accuracy has im-
034 proved steadily with deeper and wider networks such as
035 EDSR [7]. However, the accompanying growth in param-
036 eters and computation makes these models impractical for
037 resource-constrained platforms such as mobile and edge de-

vices, where memory, latency, and energy are tightly bud-
geted. This has motivated a sustained line of work on
lightweight SR, which seeks a favorable trade-off between
reconstruction quality and model cost (parameters, FLOPs,
and inference latency).

Lightweight CNN-based SR. Early efficient designs re-
duce cost through architectural economy. CARN employs
cascading residual connections with group convolutions for
a fast and compact network [2]. IMDN introduces informa-
tion multi-distillation to progressively refine and aggregate
features under a tight budget [4], and RFDN reformulates
this idea with residual feature distillation blocks [8]. RLFN
further simplifies the distillation topology to lower latency
while preserving accuracy, winning the NTIRE efficient-SR
challenge [5]. These methods are highly efficient, but their
purely convolutional, locally-connected operators have a
limited receptive field and cannot easily model the long-
range, repetitive structure that natural and graphical images
exhibit.

Lightweight Transformer-based SR. Self-attention cap-
tures long-range dependencies and has been adapted to ef-
ficient SR. SwinIR introduces window-based attention for
image restoration [6]; ELAN streamlines long-range at-
tention to cut redundant computation [13]; and SRFormer
revisits the window-channel trade-off with permuted self-
attention [14]. ATD augments local windows with an adap-
tive token dictionary to extend the effective context [12].
While effective, window attention couples the receptive
field to a fixed spatial partition, so tokens that are semanti-
cally similar but spatially distant are not directly compared
unless the window grows, which raises cost.

Content-aware / cluster-based routing. A complemen-
tary direction groups tokens by *content* rather than by spa-
tial position, so that long-range but semantically related re-
gions interact within a compact attention. SPIN performs
superpixel-based token interaction for lightweight SR [11].
CATANet pushes this idea further with a Token Aggregation
Block (TAB) that maintains a small set of routing centers,

assigns each token to its nearest center, reorders tokens by cluster label, and aggregates the reordered sequence with an intra/inter-group attention operator [9]. This content routing delivers strong accuracy at a low budget and is the most direct predecessor of our work.

Limitations of center-based routing. Despite its effectiveness, the TAB routing mechanism has four limitations. (i) Its routing centers are maintained across batches through a k -means-style iteration with an exponential moving average (EMA) buffer, so the centers are driven by a cross-batch historical state rather than strictly by the current image, and adapt only slowly to per-image semantics. (ii) The token-to-center relation is a hard $\arg\max$ label that discards the *confidence* of the assignment, treating a token that clearly belongs to one cluster identically to an ambiguous one on a boundary. (iii) The reordering clusters tokens only by label, with no intra-cluster purity constraint, so high- and low-confidence tokens are interleaved within a group during aggregation. (iv) Prototype *generation* and *confirmation* are entangled in a single iterative update, which weakens interpretability and stability. These observations suggest that a routing module built directly from the current input, that preserves and exploits assignment confidence, and that decouples prototype generation from confirmation, could improve content routing without sacrificing the lightweight regime.

Our approach. We propose *Dynamic Prototype Routing* (DPR), a drop-in replacement for the TAB routing that keeps the CATANet backbone, its external tensor interface, and its lightweight parameter budget unchanged. DPR generates a set of semantic prototypes directly from the current tokens via a soft assignment, confirms them with a learnable query bank decoupled from generation, matches tokens back to the confirmed prototypes with a learnable temperature, and propagates an explicit per-token confidence through ordering and aggregation. The resulting network, **DPRNet**, is trained on DIV2K and evaluated on five standard benchmarks at $\times 2 / \times 3 / \times 4$, where it matches or slightly surpasses the strongest content-routing baseline across benchmarks at a comparable parameter and computation budget, while exposing an interpretable, confidence-aware routing mechanism.

Contributions. Our main contributions are summarized as follows.

- **C1 – Dynamic semantic prototypes.** We generate routing prototypes from the current tokens by a soft, assignment-weighted aggregation, removing the cross-batch EMA center buffer so that prototypes follow per-image semantics (Sec. 2.3, Eqs. (4)–(5)).

- **C2 – Decoupled prototype confirmation.** We separate prototype *generation* from *confirmation* by refining the content prototypes with a learnable query bank through a single gated cross-attention, improving the stability and interpretability of the routing slots (Eq. (7)).
- **C3 – Confidence-aware routing.** We explicitly retain the token assignment confidence and use it in three places: a composite ordering key that sorts high-confidence tokens first within each cluster, a confidence-modulated global gate, and a soft prototype fallback for uncertain tokens (Eqs. (10), (11), (12)).
- **C4 – Stabilized routing.** We introduce a learnable routing temperature together with a prototype usage-balancing regularizer to prevent collapse onto a few prototypes and to reduce seed-to-seed variance (Eqs. (8), (13)).

Each design element is exposed as an independent switch that defaults to the full model and is exactly equivalent to its disabled form when turned off, so the contributions can be isolated cleanly in the ablation study (Sec. 3.4).

2. Method

2.1. Overview

We build on the lightweight super-resolution (SR) backbone of CATANet and replace its token-routing component with a *Dynamic Prototype Routing* (DPR) module. Given a low-resolution (LR) input $I_{LR} \in \mathbb{R}^{3 \times H \times W}$, a 3×3 convolution extracts shallow features $F_0 \in \mathbb{R}^{C \times H \times W}$ ($C=40$). The deep feature extractor is a stack of $L=8$ blocks; each block couples a content-routed global aggregation unit (the Token Aggregation Block, TAB) with a local self-attention unit (LRSA), followed by a 3×3 fusion convolution and a residual connection:

$$F_l = F_{l-1} + \text{Conv}_{3 \times 3}(\text{LRSA}_l(\text{TAB}_l(F_{l-1}))), \quad l = 1, \dots, L. \quad (1)$$

A global residual links the deep features to the network output, which is reconstructed by a pixel-shuffle upsampler and added to a bilinearly upsampled base image:

$$I_{SR} = \text{Up}(F_0 + F_L) + \text{Bilinear}_{\uparrow s}(I_{LR}), \quad (2)$$

where $s \in \{2, 3, 4\}$ is the scale factor. The contribution of this work is entirely inside TAB: we keep its external interface $\mathbb{R}^{C \times H \times W} \rightarrow \mathbb{R}^{C \times H \times W}$ unchanged and substitute the original means-based center iteration and hard-label sorting with the DPR pipeline (Sec. 2.3), while reusing the subsequent intra/inter-group aggregation operator (IASA, Sec. 2.4). This keeps the rest of the backbone untouched and isolates the effect of the routing mechanism.

Inside a TAB, the input feature map is flattened into a token sequence $X \in \mathbb{R}^{B \times N \times C}$ with $N=HW$, layer-normalized, routed by DPR, aggregated by IASA, and fused

by a 1×1 convolution and a convolutional feed-forward network (ConvFFN), each with a residual connection. The data flow is

$$X \rightarrow \text{DPR} \rightarrow \text{IASA} \rightarrow \text{Conv}_{1 \times 1} \rightarrow \text{ConvFFN}. \quad (3)$$

2.2. Motivation: limitations of center-based routing

The original TAB performs content-aware reordering by maintaining a set of routing centers via a k -means-style iteration with an exponential moving average (EMA) buffer, assigning each token to its nearest center by an `argmax` hard label, sorting tokens by that label, and aggregating the sorted sequence. This design has four limitations that motivate DPR: (i) the centers are driven by a cross-batch historical buffer rather than strictly by the current input, so they adapt slowly to per-image semantics; (ii) the token-to-center relation is a hard label that discards the assignment confidence; (iii) the ordering only clusters by label, with no intra-cluster purity constraint; and (iv) prototype generation and prototype confirmation are entangled, weakening interpretability. DPR addresses these points by generating prototypes from the current tokens, decoupling generation from confirmation, and propagating an explicit assignment confidence through routing and aggregation.

2.3. Dynamic Prototype Routing

DPR maps a token sequence $X \in \mathbb{R}^{B \times N \times C}$ to a content-reordered sequence together with a set of M dynamic prototypes and routing meta-information. We describe its four stages.

(1) Dynamic prototype generation. We first embed tokens with a lightweight head $E = \text{GELU}(\text{Linear}(\text{LN}(X))) \in \mathbb{R}^{B \times N \times C}$, and predict a soft assignment of each token to M prototype slots:

$$A = \text{softmax}(EW_a) \in \mathbb{R}^{B \times N \times M}, \quad W_a \in \mathbb{R}^{C \times M}, \quad (4)$$

where the softmax is taken over the M slots. Each prototype is the assignment-weighted average of the *current* tokens,

$$P_m^c = \frac{\sum_n A_{n,m} X_n}{\sum_n A_{n,m} + \epsilon}, \quad P^c = \text{normalize}(\text{LN}(P^c)), \quad (5)$$

with ℓ_2 normalization along the channel dimension. Unlike the EMA centers of the baseline, P^c depends only on the current input and carries no cross-batch state (addressing limitation (i)).

(2) Prototype confirmation by query refinement. To decouple prototype *generation* from *confirmation* (limitation (iv)), we refine P^c with a learnable query bank $Q \in \mathbb{R}^{M \times C}$ via a single cross-attention over the tokens. With

query seed $\tilde{P} = P^c + Q$, queries from the seed, and keys/values from the embedded/raw tokens,

$$P^r = \text{softmax}\left(\frac{(\tilde{P}W_q)(EW_k)^\top}{\sqrt{d_r}}\right)(XW_v), \quad (6)$$

where d_r is the router dimension. The refinement is fused through a learnable gate $\gamma = \sigma(g_r) \in (0, 1)$,

$$P = \text{normalize}(\text{LN}(P^c + \gamma P^r)). \quad (7)$$

Initializing g_r to a small value keeps the refinement a residual perturbation at the start of training, so the prototype identities remain stable across samples. This stage can be disabled (ablation A1) by setting `use_prototype_query_refine=false`, which falls back to $P = P^c$.

(3) Confidence-aware token-to-prototype matching. We match tokens *back* to the confirmed prototypes in a shared ℓ_2 -normalized router space, $\hat{e}_n = \text{normalize}(E_n W_t)$ and $\hat{p}_m = \text{normalize}(P_m W_p)$, and compute routing scores with a *learnable temperature* τ :

$$S = \text{softmax}(\tau \hat{E} \hat{P}^\top) \in \mathbb{R}^{B \times N \times M}, \quad \tau = \min(e^\theta, \tau_{\max}), \quad (8)$$

where θ is a learnable log-scale clamped by $\tau_{\max}$ for stability. A larger τ sharpens the assignment distribution and increases the separation of routing scores. For each token we read out the hard label and its confidence,

$$x_n^s = \max_m S_{n,m}, \quad b_n = \arg \max_m S_{n,m}, \quad (9)$$

and *explicitly retain* the confidence x_n^s (addressing limitation (ii)).

(4) Confidence-aware ordering. Instead of sorting only by the cluster label, we use a composite key that keeps the inter-cluster order while ordering high-confidence tokens first *within* each cluster (addressing limitation (iii)):

$$k_n = b_n + 0.5(1 - x_n^s), \quad \pi = \text{argsort}_n(k_n). \quad (10)$$

Since $x_n^s \in (0, 1)$, the additive term lies in $(0, 0.5)$ and never crosses an integer label boundary, so clusters stay contiguous while their members are purity-ordered. The token sequence, the labels, and the confidences are gathered by π into X^π , b^π , $x^{s,\pi}$, and the inverse permutation π^{-1} is recorded for scatter-back. Setting `use_conf_sort=false` reduces Eq. (10) to a pure label sort $k_n = b_n$ (ablation A2).

DPR outputs the reordered tokens X^π , the dynamic prototypes P , the full score map S , and the per-token confidence x^s , which are consumed by the aggregator below.

2.4. Confidence-aware aggregation (IASA)

The aggregation operator is retained from the backbone but is made confidence-aware. On the reordered sequence X^π it computes (a) intra-group local attention over overlapping groups of size g , capturing fine local structure within each contiguous, semantically coherent segment; and (b) a global branch that attends to the dynamic prototypes P projected to keys/values. Let O_{loc} and O_{glb} denote the two branch outputs. They are combined through a learnable global gate $\beta = \sigma(g_g)$ modulated by the per-token confidence:

$$O = O_{\text{loc}} + \beta \cdot x^{s,\pi} \cdot O_{\text{glb}}, \quad (11)$$

so that low-confidence tokens rely less on the global prototype summary. The output is scattered back to the original token order via π^{-1} and projected. Setting `use_iasa_score_gate=false` removes the confidence factor in Eq. (11) ($O = O_{\text{loc}} + \beta O_{\text{glb}}$; ablation A2).

Soft prototype fallback. To stabilize tokens whose hard assignment is uncertain, TAB adds a soft, confidence-weighted prototype context to the aggregated output. With a soft summary $\bar{X} = SP$ projected by W_s and a learnable gate $\alpha = \sigma(g_s)$,

$$Y = O + \alpha (1 - x^s) (\bar{X} W_s), \quad (12)$$

where $(1 - x^s)$ emphasizes low-confidence tokens. The result Y passes through the 1×1 convolution, the residual connection, and the ConvFFN. Setting `use_soft_fallback=false` reduces Eq. (12) to $Y = O$ (ablation A2). Equations (10), (11) and (12) together form the confidence-aware routing path, and the three switches are independent and default to the full model.

2.5. Prototype usage balancing

A learnable temperature can collapse routing onto a few prototypes. To encourage a balanced use of slots and reduce seed-to-seed variance, we add a usage-entropy regularizer. Let $u_m = \frac{1}{N} \sum_n S_{n,m}$ be the mean usage of prototype m ; its normalized entropy and the auxiliary loss are

$$\mathcal{H}(u) = -\frac{1}{\log M} \sum_m u_m \log u_m, \quad \mathcal{L}_{\text{bal}} = \lambda (1 - \mathcal{H}(u)), \quad (13)$$

which is maximized (loss minimized) when usage is uniform. The total objective adds this term, summed over all blocks, to the pixel reconstruction loss:

$$\mathcal{L} = \mathcal{L}_{\text{pix}} + \sum_l \mathcal{L}_{\text{bal}}^{(l)}. \quad (14)$$

Setting the per-block weight $\lambda=0$ disables balancing (ablation A3).

2.6. Discussion of design

DPR keeps the backbone’s interface and parameter budget in the lightweight regime while changing only how tokens are organized for aggregation. The four stages map directly to the contributions: dynamic generation (Eqs. 4–5) for C1, generation/confirmation decoupling (Eq. 7) for C2, confidence-aware routing (Eqs. 8–10, 11–12) for C3, and the learnable temperature with usage balancing (Eqs. 8, 13) for C4. Because each design element is exposed as an independent switch that defaults to the full model and is exactly equivalent to its disabled form when turned off, the contributions can be isolated cleanly in the ablation study (Sec. 3.4).

3. Experiments

3.1. Datasets and implementation details

Datasets and metrics. Following the standard lightweight-SR protocol, we train on the 800 training images of DIV2K [1] and evaluate on five benchmarks: Set5, Set14, BSD100, Urban100, and Manga109. Low-resolution inputs are produced by bicubic downscaling at $\times 2$, $\times 3$, and $\times 4$. We report PSNR and SSIM on the Y channel of the YCbCr space, discarding a border of s pixels for scale s , which matches the convention used by the cited baselines.

Network configuration. DPRNet keeps the CATANet backbone unchanged and replaces only the routing inside each TAB with the DPR module (Sec. 2). The deep feature extractor stacks $L=8$ blocks with channel width $C=40$. The number of prototype slots per block is $M = [16, 32, 64, 128, 16, 32, 64, 128]$, and the intra-group attention uses a group size of 128. The learnable routing temperature is initialized at e^6 via a log-scale and clamped at $\tau_{\text{max}}=10$ (Eq. (8)); the per-block usage-balancing weights are $\lambda = [5 \times 10^{-4}]_{\times 6}$ for the first six blocks and 7×10^{-4} for the last two (Eq. (13)). These yield 0.60–0.67 M parameters depending on scale (Table 2).

Training. We first train the $\times 2$ model from scratch for 800k iterations, and then finetune it to $\times 3$ and $\times 4$ for 250k iterations each, initializing from the $\times 2$ checkpoint. Training patches are 128×128 , 192×192 , and 256×256 HR crops for $\times 2 / \times 3 / \times 4$ respectively, with a batch size of 16 and random horizontal flips and 90° rotations for augmentation. We optimize the L1 reconstruction loss together with the usage-balancing regularizer (Eq. (13)) using Adam ($\beta_1=0.9$, $\beta_2=0.99$, no weight decay) at an initial learning rate of 2×10^{-4} , halved by a MultiStep schedule (milestones at $\{300, 500, 650, 700, 750\}k$ for $\times 2$ and $\{125, 200, 225, 237.5\}k$ for the $\times 3 / \times 4$ finetuning).

The learnable routing temperature is trained with a $0.1\times$ learning-rate multiplier for stability. Validation is run every 2.5k iterations; the unified reporting checkpoint per scale is selected by the best benchmark-averaged PSNR/SSIM (Set5–Manga109) without per-dataset cherry-picking. Experiments use PyTorch with the BasicSR framework [10] on [PLACEHOLDER: $N\times$ NVIDIA <GPU model>, total $\approx$ <GPU-hours>; fill from training machine before submission].

Efficiency measurement. Parameters, Multi-Adds (MACs reported by thop), and inference latency are measured on a single CUDA GPU for a fixed HR output of $\approx 720\times 1280$ (the LR input is derived per scale), averaging over 100 runs after 20 warm-up runs. To compare with FLOPs-based reports, the Multi-Adds values are multiplied by two.

3.2. Comparison with state-of-the-art

We compare DPRNet against eight representative lightweight SR methods spanning CNN-based designs (CARN [2], IMDN [4], RFDN [8], RLFN [5]) and Transformer/routing-based designs (SwinIR-light [6], ELAN-light [13], SRFormer-light [14], and the most direct predecessor CATANet [9]). Competitor numbers are cited verbatim from the original papers (see the notes of Table 1); DPRNet uses a single unified checkpoint per scale, reported across all five benchmarks without per-dataset selection.

Overall accuracy. Table 1 shows that DPRNet and CATANet form the top tier at all three scales, both clearly ahead of the remaining seven methods. For example, at $\times 4$ DPRNet reaches 32.59/0.9002 dB on Set5 and 27.76/0.7434 on BSD100, exceeding the strongest non-routing competitor SRFormer-light (32.51/0.8988 and 27.73/0.7422) while using a comparable parameter budget. The consistent advantage of the two content-routing methods over window- and distillation-based baselines supports the value of content-aware token aggregation in the lightweight regime.

Comparison with CATANet. Against CATANet—the strongest baseline and the method DPRNet differs from only in the routing module—DPRNet attains the best result on nearly every benchmark/scale cell, with CATANet consistently second. The margins are small but consistent: at $\times 4$ DPRNet improves Urban100 from 26.87/0.8081 to 26.89/0.8089, Set5 from 32.58/0.8998 to 32.59/0.9002, Set14 from 28.90/0.7880 to 28.91/0.7883, and Manga109 PSNR from 31.31 to 31.32; at $\times 2$ it reaches the best Set14 (34.06/0.9220), Urban100 PSNR (33.15), and Manga109 (39.38/0.9789); and at $\times 3$ it leads or ties CATANet on every benchmark (e.g. BSD100 29.29/0.8103 and Urban100

29.06/0.8689). The single metric on which CATANet remains ahead is Manga109 SSIM at $\times 4$ (0.9183 vs. 0.9181). These gaps are within a few hundredths of a dB, so the overall message is that replacing the EMA-center routing of TAB with the dynamic prototype routing of DPR matches and slightly surpasses CATANet-level reconstruction quality while changing the routing mechanism to the interpretable, confidence-aware formulation analyzed in Sec. 3.4 and Sec. 3.5.

Efficiency. Table 2 reports the efficiency profile at $\times 4$. DPRNet has 660K parameters and 45.77 G Multi-Adds for a fixed $\approx 720\times 1280$ output, keeping it within the lightweight bracket alongside the compared methods while delivering top-tier accuracy. DPRNet’s parameter count is modestly higher than CATANet’s (660K vs. 535K at $\times 4$), reflecting the learnable query bank and router projections of DPR; the per-scale profile is given in Table 3. A complete Multi-Adds/latency comparison against all competitors under our measurement protocol is left for the final version, as those numbers are not yet collected under a unified setting.

3.3. Model complexity and efficiency

Table 2 compares model complexity at $\times 4$, and Table 3 gives DPRNet’s per-scale profile. All DPRNet numbers are measured under the protocol of Sec. 3.1 (fixed $\approx 720\times 1280$ output, thop MACs, latency over 100 runs after 20 warm-ups). For CATANet we additionally cite its reported $\times 4$ Multi-Adds from the original paper, measured at the same output size and therefore directly comparable to our MACs.

Parameters. DPRNet stays within the lightweight bracket at 660K parameters ($\times 4$), comparable to Transformer baselines such as ELAN-light (601K) and well below SwinIR-light (897K), SRFormer-light (873K), and the CNN-heavy CARN (1592K). It is, however, larger than its direct predecessor CATANet (535K): the learnable query bank and the router/score projections of DPR add parameters relative to the EMA-center routing of CATANet. We therefore do not claim a parameter advantage over CATANet; the comparison below is at *matched-or-better accuracy* (Urban100 26.89 vs. 26.87, with the full ablation analysis in Sec. 3.4).

Computation. In terms of Multi-Adds, DPRNet requires 45.77 G at $\times 4$, which is in fact slightly *lower* than CATANet’s 46.8 G at the same output resolution. Thus the extra parameters of DPR do not translate into extra inference computation: the dynamic prototype routing replaces CATANet’s training-time EMA center updates with per-image query refinement that is gated by a small

Table 1. Quantitative comparison with state-of-the-art lightweight SR methods on five benchmarks (PSNR(dB)/SSIM on the Y channel). **Best** and second-best results are highlighted. Competitor metrics are cited from the original papers (see notes). “-” denotes results not reported in the source.

Method	Params (K)	Set5		Set14		BSD100		Urban100		Manga109	
		PSNR	SSIM	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM
Scale ×2											
CARN	1592	37.76	0.9590	33.52	0.9166	32.09	0.8978	31.92	0.9256	38.36	0.9765
IMDN	694	38.00	0.9605	33.63	0.9177	32.19	0.8996	32.17	0.9283	38.88	0.9774
RFDN	534	38.05	0.9606	33.68	0.9184	32.16	0.8994	32.12	0.9278	38.88	0.9773
RLFN	527	38.07	0.9607	33.72	0.9187	32.22	0.9000	32.33	0.9299	-	-
SwinIR-light	878	38.14	0.9611	33.86	0.9206	32.31	0.9012	32.76	0.9340	39.12	0.9783
ELAN-light	582	38.17	0.9611	33.94	0.9207	32.30	0.9012	32.76	0.9340	39.11	0.9782
SRFormer-light	853	38.23	<u>0.9613</u>	33.94	0.9209	32.36	0.9019	32.91	<u>0.9353</u>	39.28	<u>0.9785</u>
CATANet	477	<u>38.28</u>	0.9617	<u>33.99</u>	<u>0.9217</u>	<u>32.37</u>	<u>0.9023</u>	<u>33.09</u>	0.9372	<u>39.37</u>	0.9784
DPRNet (ours)	602	38.29	0.9617	34.06	0.9220	32.38	0.9031	33.15	0.9372	39.38	0.9789
Scale ×3											
CARN	1592	34.29	0.9255	30.29	0.8407	29.06	0.8034	28.06	0.8493	33.43	0.9427
IMDN	703	34.36	0.9270	30.32	0.8417	29.09	0.8046	28.17	0.8519	33.61	0.9445
RFDN	541	34.41	0.9273	30.34	0.8420	29.09	0.8050	28.21	0.8525	33.67	0.9449
RLFN	-	-	-	-	-	-	-	-	-	-	-
SwinIR-light	886	34.62	0.9289	30.54	0.8463	29.20	0.8082	28.66	0.8624	33.98	0.9478
ELAN-light	590	34.64	0.9288	30.55	0.8463	29.21	0.8081	28.69	0.8624	34.00	0.9478
SRFormer-light	861	34.67	0.9296	<u>30.57</u>	<u>0.8469</u>	29.26	0.8099	28.81	<u>0.8655</u>	34.19	<u>0.9489</u>
CATANet	550	<u>34.75</u>	0.9300	30.67	0.8481	<u>29.28</u>	<u>0.8101</u>	<u>29.04</u>	0.8689	<u>34.40</u>	0.9499
DPRNet (ours)	674	34.76	0.9301	30.67	0.8481	29.29	0.8103	29.06	0.8689	34.41	0.9499
Scale ×4											
CARN	1592	32.13	0.8937	28.60	0.7806	27.58	0.7349	26.07	0.7837	30.42	0.9070
IMDN	715	32.21	0.8948	28.58	0.7811	27.56	0.7353	26.04	0.7838	30.45	0.9075
RFDN	550	32.24	0.8952	28.61	0.7819	27.57	0.7360	26.11	0.7858	30.58	0.9089
RLFN	543	32.24	0.8952	28.62	0.7813	27.60	0.7364	26.17	0.7877	-	-
SwinIR-light	897	32.44	0.8976	28.77	0.7858	27.69	0.7406	26.47	0.7980	30.92	0.9151
ELAN-light	601	32.43	0.8975	28.78	0.7858	27.69	0.7406	26.54	0.7982	30.92	0.9150
SRFormer-light	873	32.51	0.8988	28.82	0.7872	27.73	0.7422	26.67	0.8032	31.17	0.9165
CATANet	535	<u>32.58</u>	<u>0.8998</u>	<u>28.90</u>	<u>0.7880</u>	<u>27.75</u>	<u>0.7427</u>	<u>26.87</u>	<u>0.8081</u>	<u>31.31</u>	0.9183
DPRNet (ours)	660	32.59	0.9002	28.91	0.7883	27.76	0.7434	26.89	0.8089	31.32	<u>0.9181</u>

Notes. DPRNet (ours) uses a single unified checkpoint per scale (x2: net_g_792500; x3/x4: net_g_250000), reported across all benchmarks without per-dataset cherry-picking. Competitor numbers are cited verbatim: CARN/IMDN/RFDN/SwinIR-light/ELAN-light/CATANet from CATANet (CVPR 2025, Tab. 2); SRFormer-light from its paper (Tab. VI); RLFN from its paper (Tab. 1). RLFN does not report x3 or Manga109 in the original paper (“-”).

Params for SwinIR-light and ELAN-light follow the convention of the CATANet table.

444 scalar (Eq. (7)), keeping the dominant cost—the intra-
445 group attention—essentially unchanged. DPRNet therefore
446 matches or slightly surpasses CATANet-level reconstruc-
447 tion at comparable compute, trading a modest parameter in-
448 crease for an interpretable, confidence-aware routing mech-
449 anism.

450 **Per-scale behavior.** Table 3 shows the expected trend
451 across scales: for a fixed HR output the LR feature

map shrinks as the scale grows, so Multi-Adds drop
from 126.52 G (×2) to 64.28 G (×3) and 45.77 G (×4),
with latency following the same order (712.16, 285.39,
196.77 ms). The parameter count is essentially scale-
independent (≈ 0.60 – 0.67 M), as only the input projection
and the upsampling tail change with scale.

Pending measurements. A complete Multi-Adds and la-
tency comparison against all eight competitors under a

Table 2. Efficiency comparison on $\times 4$ SR. Multi-Adds are reported by thop (MACs); to compare against FLOPs-based papers, multiply by 2. “-” denotes values not yet collected under our protocol. Representative PSNR is on Urban100 ($\times 4$).

Method	Params (K)	Multi-Adds (G)	Latency (ms)	Urban100 PSNR
CARN	1592	-	-	26.07
IMDN	715	-	-	26.04
RFDN	550	-	-	26.11
RLFN	543	-	-	26.17
SwinIR-light	897	60.3 [‡]	-	26.47
ELAN-light	601	-	-	26.54
SRFormer-light	873	56.5 [‡]	-	26.67
CATANet	535	46.8 [‡]	-	26.87
DPRNet (ours)	660	45.77	196.77 $\pm$ 10.02	26.89

Notes. DPRNet efficiency is measured on a single CUDA GPU with a fixed HR output of $\approx 720 \times 1280$ (LR input derived per scale), warmup=20, repeat=100. Multi-Adds are MACs reported by thop.

[‡]Multi-Adds for CATANet, SwinIR-light and SRFormer-light are cited from the CATANet paper (CVPR’25, Tab. 6), measured with a $3 \times 256 \times 256$ input ($\times 4$ output 1024×1024); this convention differs from our $\approx 720 \times 1280$ setting, so these values are *not* directly comparable to DPRNet and serve only as a reference. Remaining competitors’ Multi-Adds/latency are not yet collected under a unified protocol (“-”); fill before submission.

Table 3. DPRNet per-scale efficiency profile (measured). LR input derived from a fixed HR output $\approx 720 \times 1280$; Multi-Adds are MACs (thop); latency averaged over 100 runs after 20 warmups on a single CUDA GPU.

Scale	LR Input	Params (K)	Multi-Adds (G)	Latency (ms)
$\times 2$	360 $\times$ 640	601.95	126.52	712.16 $\pm$ 23.47
$\times 3$	240 $\times$ 426	674.15	64.28	285.39 $\pm$ 5.66
$\times 4$	180 $\times$ 320	659.71	45.77	196.77 $\pm$ 10.02

single unified protocol is not yet available: published competitor compute figures follow inconsistent parameter and measurement conventions, and latency is hardware-dependent and cannot be borrowed from other papers’ devices. We report the quantities we can state consistently here—DPRNet’s full measured profile and the directly comparable CATANet Multi-Adds—and defer the unified competitor table to the final version (re-measured under our protocol or taken from a single common source). We do not fabricate the missing entries, which appear as “-” in Table 2.

3.4. Ablation study

We ablate the four design elements of DPR using the independent switches exposed in Sec. 2. All ablations are conducted at scale $\times 4$, following CATANet [9] whose ablations are likewise performed at $\times 4$, the setting in which content routing on high-frequency texture is most exercised. To keep the compute budget tractable, each variant is *fine-tuned* from the converged full-model checkpoint and trained for a short 80k-iteration budget, with validation on Set5

and Urban100 (Y channel, matched `crop_border`) every 10k iterations; we report the iter-80k point. Besides PSNR/SSIM we log routing diagnostics per TAB block—the normalized prototype-usage entropy $\mathcal{H}(u)$ of Eq. (13), the active-prototype usage, and the learnable temperature—to inspect the routing behavior directly.

Protocol caveat. Because every variant is finetuned from the *same* full-model initialization for a short budget, the disabled components are removed from an already converged representation rather than learned from scratch. Under this protocol the full model attains the best reconstruction in each ablation group, but the margins are small (up to ~ 0.06 – 0.08 dB, against an ~ 0.03 dB run-to-run fluctuation in the full model’s own validation PSNR between adjacent 10k checkpoints) and the intermediate configurations are not strictly monotone in the number of enabled switches. We therefore read the PSNR/SSIM columns as evidence that each design element is *quality-neutral-to-mildly-positive and component-compatible*—the full configuration is never worse and is consistently the best, while no single switch can be credited with an isolated gain under this protocol—and we substantiate the routing-level effects through the diagnostics. A fully separated attribution of each switch would require retraining every variant from a switch-neutral initialization to the CATANet $\times 4$ from-scratch budget, which we leave to future work.

A1: decoupled prototype confirmation. Table 4 compares DPR with and without the query-refinement stage (Eq. (7)). Enabling refine gives the better result on both benchmarks (Set5 32.59 vs. 32.55 dB; Urban100 26.89 vs. 26.84 dB), confirming that adding the learnable query bank improves—and certainly does not perturb—reconstruction. As the refinement is gated by a small-initialized scalar (Eq. (7)), it acts as a residual confirmation step that keeps the prototype slots stable while decoupling generation from confirmation (C2); we therefore retain it by default. The motivation for dynamic, per-image prototypes over the cross-batch EMA centers of the baseline (C1) is supported indirectly by the published CATANet results rather than by this controlled switch, since re-implementing the EMA center branch is outside the scope of our ablation.

A2: confidence-aware routing. Table 5 adds the three confidence-aware components cumulatively: confidence-aware ordering (Eq. (10)), the IASA score gate (Eq. (11)), and the soft prototype fallback (Eq. (12)). The full configuration (v4) is the best of the four on both Set5 and Urban100, but the intermediate steps are not monotone—adding the ordering and gate alone slightly lowers PSNR before the complete set recovers and surpasses the hard-routing baseline. Given the shared-initialization protocol

above, this indicates that the three switches are mutually compatible and that the complete confidence-aware path is quality-preserving-to-mildly-positive, rather than each switch being individually accuracy-additive on top of a converged backbone. The components are kept because they make the routing path interpretable and confidence-driven: the retained per-token confidence x^s modulates ordering, the global gate, and the fallback (C3), which we visualize through the x^s distributions in Sec. 3.5.

A3: prototype usage balancing. Table 6 isolates the usage-balancing regularizer (Eq. (13)). Enabling it raises the block-averaged normalized usage entropy from 0.6227 to 0.6865 on Urban100, and the increase is *consistent across all eight TAB blocks* (b_0 – b_7 ; Table 7), i.e. the regularizer drives a more balanced use of prototype slots and counteracts the collapse that the learnable temperature would otherwise encourage (C4). The reconstruction quality is also slightly higher with the regularizer enabled (Set5 32.59 vs. 32.57 dB; Urban100 26.89 vs. 26.83 dB), so the balancing does not trade accuracy for a flatter usage distribution. The entropy effect is visible even under the short finetune protocol because it concerns the usage *distribution* rather than absolute reconstruction. The complementary part of C4—a reduction of seed-to-seed variance—requires repeating the runs over multiple seeds; with a single seed available we report the balancing mechanism here and leave the variance analysis as future work, and we do not claim a variance reduction without those runs.

Summary. The ablations show that DPR’s design elements are quality-neutral-to-mildly-positive and mutually compatible under a shared-initialization finetune—the full model is the best variant in every group—and that the balancing regularizer produces a measurable, block-consistent increase in prototype-usage entropy. Stronger, fully separated accuracy attributions are deferred to from-scratch re-training; the present study supports the mechanism of each component without overstating its isolated contribution to PSNR.

3.5. Routing interpretability analysis

To inspect what DPR actually does at routing time—rather than its effect on PSNR—we analyze the routing diagnostics logged during validation at $\times 4$ (iter 80k, Urban100). For each of the $L=8$ TAB blocks we record the per-token routing confidence $x^s = \max_m S_{n,m}$ (Eq. (9)), the number of *active* prototypes (slots receiving at least one token), the normalized usage entropy $\mathcal{H}(u)$ (Eq. (13)), and the learned temperature τ (Eq. (8)). The blocks use $M = [16, 32, 64, 128, 16, 32, 64, 128]$ prototypes respectively.

Table 4. Ablation A1 ($\times 4$): prototype query refine. PSNR(dB)/SSIM (Y channel) at iter 80k and block-averaged normalized prototype-usage entropy on Urban100. Enabling refine gives the better reconstruction on both benchmarks under the short finetune protocol; it is retained as it improves (and does not degrade) quality while keeping slot assignment stable.

Variant	Set5		Urban100
	PSNR	SSIM	PSNR/SSIM
w/o refine	32.5494	0.8995	26.8358 / 0.8065
w/ refine (default)	32.5878	0.9001	26.8905 / 0.8086

Notes. $\times 4$ ablation finetuned from the full-functioning checkpoint (net_g_250000) for an 80k budget; metrics read at iter 80k. Enabling refine yields the better PSNR/SSIM on both Set5 and Urban100, though all variants share this initialization and a short budget so the margin (~ 0.04 – 0.05 dB) should not be over-read as a large isolated gain. The design motivation of dynamic prototypes vs. history-center+EMA is supported indirectly by the cited CATANet results (Table I), not by this controlled ablation.

Table 5. Ablation A2 ($\times 4$): incremental confidence-aware routing components. PSNR(dB)/SSIM (Y channel) at iter 80k. The full configuration (v4) attains the best reconstruction, but the intermediate steps are not monotone; under the short finetune protocol the components are mutually compatible and quality-preserving rather than individually accuracy-additive.

Variant	Switches			Set5		Urban100
	sort	gate	fb	PSNR	SSIM	PSNR/SSIM
v1 (hard sort)	✗	✗	✗	32.5459	0.8995	26.8559 / 0.8074
v2 (+conf. sort)	✓	✗	✗	32.5249	0.8995	26.8324 / 0.8065
v3 (+score gate)	✓	✓	✗	32.5116	0.8992	26.8123 / 0.8056
v4 (full, +soft fb)	✓	✓	✓	32.5878	0.9001	26.8905 / 0.8086

Notes. “sort”=confidence-aware sort key, “gate”=IASA score gating, “fb”=soft fallback. All variants are $\times 4$ finetuned from the full-functioning checkpoint (net_g_250000) for an 80k budget; metrics read at iter 80k. The full configuration (v4) is the best of the four, but because the variants share this initialization and a short budget the intermediate PSNR/SSIM steps are not monotone; the components are reported as mutually compatible and quality-preserving rather than as a monotone incremental gain.

Balanced prototype usage (C4). Fig. 1 reports the normalized usage entropy per block with and without the balancing regularizer. With balancing enabled (the full model) the entropy is higher in *all eight* blocks, raising the block-mean from 0.6227 to 0.6865; the effect is largest in the high-frequency-oriented block b_4 (0.7088 $\rightarrow$ 0.8466). Consistently, the average number of active prototypes increases from 47.5 to 52.3 across blocks, i.e. fewer slots are left unused. Because a learnable temperature can otherwise sharpen the assignment distribution and concentrate tokens on a handful of prototypes, these two measurements together show that the regularizer keeps prototype usage

Table 6. Ablation A3 ($\times 4$): prototype balance loss. Block-averaged normalized prototype-usage entropy on Urban100 (higher = more balanced usage) and reconstruction quality at iter 80k. The balance loss yields consistently higher usage entropy across all eight TAB blocks while also slightly improving reconstruction quality.

Variant	Usage entropy (mean, $\uparrow$)	Set5		Urban100 PSNR/SSIM
		PSNR	SSIM	
w/o balance loss	0.6227	32.5689	0.8997	26.8262 / 0.8062
w/ balance loss (def.)	0.6865	32.5878	0.9001	26.8905 / 0.8086

Table 7. Per-block normalized prototype-usage entropy (Urban100, iter 80k). The balance loss increases usage entropy in every TAB block (b_0 – b_7).

Variant	b_0	b_1	b_2	b_3	b_4	b_5	b_6	b_7
w/o balance	0.6482	0.5655	0.6947	0.6268	0.7088	0.5764	0.6279	0.5332
w/ balance	0.7181	0.5845	0.7227	0.6810	0.8466	0.6352	0.6848	0.6193

Notes. $\times 4$ ablation finetuned from net_g.250000 for an 80k budget; diagnostics read at iter 80k on Urban100. The balance loss consistently increases prototype-usage entropy (more balanced usage) across all blocks, supporting C4’s balancing mechanism, and the PSNR/SSIM is slightly higher with it enabled. The seed-variance reduction sub-claim of C4 requires multi-seed runs (currently a single seed, 3407) and is left as future work; it must not be asserted before those runs exist.

spread out and prevents slot collapse, which is exactly the stabilizing role claimed for C4.

Bounded routing temperature (C4). The learnable temperature is clamped by $\tau_{\max}=10$ for stability (Eq. (8)). In the trained full model it settles in a moderate range across blocks (per-block τ from 5.65 to 6.51, block-mean 6.06), staying well below the clamp rather than saturating it. The temperature therefore sharpens the routing scores enough to separate prototypes while the usage regularizer keeps the assignment from collapsing—the two C4 components act as intended without hitting their safety bound.

Confidence-aware routing signal (C3). The per-token confidence x^s that drives the ordering key (Eq. (10)), the IASA gate (Eq. (11)) and the soft fallback (Eq. (12)) is logged as a non-degenerate signal in every block (block-mean x^s on the order of 10^{-2} , well above the uniform-assignment floor $1/M$ for the large- M blocks), confirming that the softmax assignment is informative rather than flat and that the confidence used by the C3 components is meaningful. A finer-grained, per-token view of this signal (the x^s histogram before vs. after the learnable temperature, and the routing cluster maps obtained by reshaping the original-order labels b_n to $H \times W$) requires a forward pass over the test images; the model already caches the original-order routing map and confidences for this purpose (`last_routing_map`, `last_x_scores`), and we




Figure 1. Normalized prototype-usage entropy $\mathcal{H}(u)$ (Eq. (13); higher = more balanced) per TAB block (b_0 – b_7) at $\times 4$, iter 80k on Urban100. The usage-balancing regularizer increases the entropy in *all eight* blocks (block-mean 0.6227 $\rightarrow$ 0.6865), confirming that it counteracts the collapse onto a few prototypes encouraged by the learnable temperature (C4).

present these qualitative visualizations together with the visual comparison in the camera-ready, keeping the quantitative interpretability claims in this section grounded in the logged statistics.

Summary. The logged diagnostics give direct, mechanism-level evidence for the routing design: the balancing regularizer produces a block-consistent increase in usage entropy and active-prototype count (C4), the learnable temperature stays in a stable, unsaturated range (C4), and the retained confidence signal is informative across all blocks (C3). These observations characterize the behavior of DPR independently of the reconstruction metrics discussed in Sec. 3.4.

4. Discussion

Content routing vs. spatial locality. The benchmark results in Sec. 3.2 indicate where content-aware routing helps most. On Urban100 and Manga109—images dominated by repetitive structures and long-range self-similarity—DPRNet and CATANet, the two routing methods, clearly separate from the window- and distillation-based baselines, whereas on the more stochastic textures of BSD100 the margin between all top-tier methods narrows. This is consistent with the premise of content routing: grouping tokens by semantics is most beneficial when semantically related content is spatially scattered, and less so when local statistics already capture the signal. DPR does not discard spatial locality but complements it, since each TAB is paired with a local self-attention unit (Sec. 2.1); the routing module supplies the long-range, content-based interaction that a fixed spatial window cannot.

The risk of over-strong routing. A confidence-aware router can, in principle, sharpen its assignments until tokens collapse onto a few prototypes, which would defeat the purpose of routing and destabilize training. Two components of DPR are explicitly designed against this failure mode: the routing temperature is clamped (Eq. (8)) and the usage-balancing regularizer (Eq. (13)) penalizes peaky usage distributions. The diagnostics in Sec. 3.5 show both mechanisms operating as intended—the learned temperature settles in a moderate, unsaturated range (block-mean 6.06, below the clamp of 10) and the balancing regularizer raises the prototype-usage entropy in all eight blocks ($0.6227 \rightarrow 0.6865$) while increasing the number of active prototypes. In other words, the gains of confidence-aware routing are obtained without letting the router become degenerate.

Prototype count and the cost trade-off. DPR uses a per-block prototype budget $M = [16, 32, 64, 128, 16, 32, 64, 128]$, growing within each stage. More prototypes give the router a finer content vocabulary but enlarge the query bank and the matching cost, while too few prototypes force unrelated tokens to share a slot. We keep the prototype schedule aligned with the CATANet backbone so that the comparison isolates the routing mechanism rather than a capacity change; this is also why DPRNet carries a modestly higher parameter count than CATANet (660K vs. 535K at $\times 4$) yet a slightly lower Multi-Adds (45.77 G vs. 46.8 G), as the added parameters live in the query bank and projections rather than in the dominant attention cost (Sec. 3.3). A systematic sweep of the prototype schedule is a natural extension but is orthogonal to the routing formulation studied here.

Limitations. Several limitations bound the present study and frame our claims conservatively. First, the ablations (Sec. 3.4) finetune every variant from the same converged full-model checkpoint under a short 80k-iteration, single-seed budget; this isolates whether each design element is quality-neutral-to-mildly-positive and mutually compatible, but it cannot attribute an isolated PSNR gain to any single switch, and the seed-variance component of C4 is therefore left as future work rather than asserted. A fully separated attribution would require retraining each variant from a switch-neutral initialization at the CATANet $\times 4$ from-scratch budget. Second, the motivation for dynamic, per-image prototypes over cross-batch EMA centers (C1) is supported indirectly through the published CATANet results rather than through a re-implemented EMA branch. Third, the unified efficiency comparison is incomplete: competitor Multi-Adds and latency are not yet measured under a single protocol, and the corresponding cells are left blank rather than filled from inconsistent sources (Sec. 3.3). Fi-

nally, models are trained on DIV2K only; larger training sets (e.g. DF2K) and real-world or arbitrary-scale degradations are not explored here and may further shift the trade-off.

5. Conclusion

We presented Dynamic Prototype Routing (DPR), a drop-in replacement for the center-based Token Aggregation Block of CATANet that keeps the backbone, the tensor interface, and the lightweight budget unchanged. DPR generates semantic prototypes directly from the current tokens, decouples their confirmation through a learnable query bank, matches tokens back with a learnable temperature, and propagates an explicit per-token confidence through ordering, gating, and a soft fallback, with a usage-balancing regularizer that keeps prototype use spread out. On five standard benchmarks at $\times 2 / \times 3 / \times 4$, the resulting DPRNet matches or slightly surpasses the strongest content-routing baseline at a comparable parameter and computation budget, while the routing diagnostics give mechanism-level evidence that the temperature stays unsaturated and the prototype usage becomes measurably more balanced across all blocks. Because DPR preserves the external interface of the routing block, it is not tied to this particular backbone and can serve as a generic, confidence-aware content-routing module for other token-aggregation architectures. Future work includes separating the contribution of each design element through from-scratch retraining from a switch-neutral initialization and a multi-seed variance study, completing the unified efficiency comparison under a single measurement protocol, and extending DPR to larger training sets and real-world or arbitrary-scale degradations.

References

- [1] Eirikur Agustsson and Radu Timofte. NTIRE 2017 challenge on single image super-resolution: Dataset and study. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 126–135, 2017.
- [2] Namhyuk Ahn, Byungkong Kang, and Kyung-Ah Sohn. Fast, accurate, and lightweight super-resolution with cascading residual network. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 252–268, 2018. arXiv:1803.08664.
- [3] Chao Dong, Chen Change Loy, Kaiming He, and Xiaoou Tang. Learning a deep convolutional network for image super-resolution. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 184–199, 2014.
- [4] Zheng Hui, Xinbo Gao, Yunchu Yang, and Xiumei Wang. Lightweight image super-resolution with information multi-distillation network. In *Proceedings of the 27th ACM International Conference on Multimedia (ACM MM)*, pages 2024–2032, 2019. arXiv:1909.11856.

- [5] Fangyuan Kong, Mingxi Li, Songwei Liu, Ding Liu, Jingwen He, Yang Bai, Fangmin Chen, and Lean Fu. Residual local feature network for efficient super-resolution. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 766–776, 2022. arXiv:2205.07514; NTIRE 2022 Efficient SR Challenge winner.
- [6] Jingyun Liang, Jiezhong Cao, Guolei Sun, Kai Zhang, Luc Van Gool, and Radu Timofte. SwinIR: Image restoration using swin transformer. In *Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops (ICCVW)*, pages 1833–1844, 2021. arXiv:2108.10257.
- [7] Bee Lim, Sanghyun Son, Heewon Kim, Seungjun Nah, and Kyoung Mu Lee. Enhanced deep residual networks for single image super-resolution. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 136–144, 2017. arXiv:1707.02921.
- [8] Jie Liu, Jie Tang, and Gangshan Wu. Residual feature distillation network for lightweight image super-resolution. In *Proceedings of the European Conference on Computer Vision Workshops (ECCVW)*, 2020. arXiv:2009.11551; AIM 2020 Efficient SR Challenge.
- [9] Xin Liu, Jie Liu, Jie Tang, and Gangshan Wu. CATANet: Efficient content-aware token aggregation for lightweight image super-resolution. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17902–17912, 2025. arXiv:2503.06896.
- [10] Xintao Wang, Liangbin Xie, Ke Yu, Kelvin C.K. Chan, Chen Change Loy, and Chao Dong. BasicSR: Open source image and video restoration toolbox. <https://github.com/XPixelGroup/BasicSR>, 2020.
- [11] Aiping Zhang, Wenqi Ren, Yi Liu, and Xiaochun Cao. Lightweight image super-resolution with superpixel token interaction. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- [12] Leheng Zhang, Yuxuan Li, Xiaorui Zhou, Xiaojie Zhao, and Shuhang Gu. Transcending the limit of local window: Advanced super-resolution transformer with adaptive token dictionary. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. ATD; arXiv:2401.08209.
- [13] Xindong Zhang, Hui Zeng, Shi Guo, and Lei Zhang. Efficient long-range attention network for image super-resolution. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 649–667, 2022. arXiv:2203.06697.
- [14] Yupeng Zhou, Zhen Li, Chun-Le Guo, Song Bai, Ming-Ming Cheng, and Qibin Hou. SRFormer: Permuted self-attention for single image super-resolution. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 12780–12791, 2023. arXiv:2303.09735.