

# Dynamic Prototype Routing for Lightweight Image Super-Resolution

Author Name <sup>1,\*</sup>

<sup>1</sup> Department, Institution, City, Country; author@example.edu

\* Correspondence: author@example.edu

## Featured Application

Confidence-aware content routing for efficient single-image super-resolution on resource-constrained platforms such as mobile and edge devices.

## Abstract

Lightweight single image super-resolution (SR) must balance reconstruction quality with model cost on resource-constrained platforms. Content-aware routing improves this balance by grouping semantically related tokens, but the Token Aggregation Block (TAB) of CATANet uses cross-batch exponential-moving-average centers, hard assignment labels, and a coupled prototype update/confirmation step. We propose *Dynamic Prototype Routing* (DPR), a drop-in TAB replacement that keeps the CATANet backbone and tensor interface unchanged. DPR forms per-image prototypes from the current tokens, refines them with a learnable query bank, matches tokens with a bounded learnable temperature, and carries assignment confidence into token ordering, global gating, and a soft fallback; a balancing loss discourages prototype-slot collapse. Trained on DIV2K and evaluated on five standard benchmarks at  $\times 2 / \times 3 / \times 4$ , DPRNet stays in the lightweight regime and matches CATANet-level accuracy. It yields small gains on most benchmark cells, including 26.89 dB on Urban100 at  $\times 4$  with 52.14 G Multi-Adds. It uses fewer Multi-Adds than the lightweight Transformer baselines SwinIR-light and SRFormer-light under the same input convention. Multi-seed ablations confirm that the dynamic prototypes outperform a history-center variant with an exponential moving average and that the balancing loss makes prototype usage both more balanced and more stable across seeds, while the routing diagnostics remain interpretable with soft assignment confidence.

**Keywords:** image super-resolution; lightweight super-resolution; content-aware routing; dynamic prototype routing; self-attention; vision transformer; deep learning

## 1. Introduction

Single image super-resolution (SR) aims to reconstruct a high-resolution (HR) image from a low-resolution (LR) observation, a problem that underpins applications from mobile photography to display rendering. Before deep learning, SR was often formulated with image priors such as sparse representation and internal self-similarity [1,2]. Since the seminal deep model SRCNN [3], accuracy has improved through deeper, recurrent, progressive, and feature-reuse networks, including VDSR [4], DRCN [5], LapSRN [6], DBPN [7], RDN [8], RCAN [9], and EDSR [10]. A recent survey summarizes this progression across architectures, losses, datasets, and metrics [11]. The corresponding growth in parameters

Received:

Revised:

Accepted:

Published:

**Copyright:** © 2026 by the authors.

Submitted to *Appl. Sci.* for possible

open access publication under the

terms and conditions of the [Creative Commons Attribution \(CC BY\) license](#).

[Commons Attribution \(CC BY\) license](#).

and computation, however, makes many high-capacity models impractical for mobile and edge devices, where memory, latency, and energy are tightly constrained. This tension has motivated a sustained line of *lightweight* SR, which seeks a favorable trade-off between reconstruction quality and model cost.

#### Lightweight CNN-based SR.

Efficient CNN designs reduce cost by moving computation to low-resolution features, reusing parameters, simplifying feature distillation, or replacing heavy operators with compact mixing. FSRCNN accelerates SRCNN-style reconstruction [12], ESPCN performs efficient sub-pixel upsampling [13], and recursive or memory-based designs reuse features across depth [14,15]. CARN employs cascading residual connections with group convolutions [16]; IMDN introduces information multi-distillation [17]; RFDN reformulates this idea with residual feature distillation blocks [18]; and RLFN further simplifies local feature extraction for efficient SR [19]. More recent lightweight designs also explore shuffle-based convolutional mixing and dual-regression constraints [20,21]. Efficient-SR challenges reflect the same deployment pressure by ranking models with accuracy, runtime, parameters, and FLOPs considered together [22]. For the present work, the main limitation of these methods is not efficiency but context: locally connected convolutional operators cannot directly compare spatially distant yet semantically similar regions without stacking many layers.

#### Lightweight Transformer-based SR.

Long-range modeling tackles the context limitation from a different direction. Non-local operators aggregate features across distant positions [23], and Transformer-based restoration extends this idea with learned attention. IPT shows that large pre-trained Transformers can serve image restoration tasks including SR [24], while Swin Transformer provides the shifted-window operator that later restoration models adopt for tractable attention [25]. SwinIR introduces window-based attention for image restoration [26]; ELAN streamlines long-range attention to cut redundant computation [27]; SRFormer revisits the window-channel trade-off with permuted self-attention [28]; and HAT activates more pixels to enlarge the effective receptive field [29]. Efficient variants further mix local operators with omni-dimensional, n-gram, or hierarchical attention patterns [30–33]. ATD augments local windows with an adaptive token dictionary to extend the effective context [34]. While effective, window attention still ties token interaction to a fixed spatial partition. Semantically similar but spatially distant tokens therefore interact only when the window grows or an additional global path is introduced.

#### Content-aware / cluster-based routing.

A complementary direction groups tokens by *content* rather than by spatial position, allowing long-range but semantically related regions to interact within a compact attention. This idea is related to earlier self-exemplar and non-local SR designs, where similar patches or texture correspondences are explicitly mined across an image or reference [2,35,36]. SPIN performs superpixel-based token interaction for lightweight SR [37]. CATANet pushes this idea further with a Token Aggregation Block (TAB) that maintains a small set of routing centers, assigns each token to its nearest center, reorders tokens by cluster label, and aggregates the reordered sequence with an intra/inter-group attention operator [38]. This content routing delivers strong accuracy at a low budget and is the most direct predecessor of our work.

## Limitations of center-based routing.

Despite its effectiveness, the TAB routing mechanism has four limitations. (i) Its routing centers are maintained across batches through a  $k$ -means-style iteration with an exponential moving average (EMA) buffer, so the centers are driven by a cross-batch historical state rather than strictly by the current image and adapt only slowly to per-image semantics. (ii) The token-to-center relation is a hard  $\text{argmax}$  label that discards the *confidence* of the assignment, treating a token that clearly belongs to one cluster identically to an ambiguous one on a boundary. (iii) The reordering clusters tokens only by label, with no intra-cluster purity constraint, so high- and low-confidence tokens are interleaved within a group during aggregation. (iv) Prototype *generation* and *confirmation* are entangled in a single iterative update, which weakens interpretability and stability. These observations suggest that a routing module built directly from the current input, that preserves and exploits assignment confidence, and that separates prototype generation from confirmation could improve content routing without sacrificing the lightweight regime.

## Our approach.

We propose *Dynamic Prototype Routing* (DPR), a drop-in replacement for the TAB routing that keeps the CATANet backbone and its external tensor interface unchanged while remaining in a comparable lightweight operating regime. DPR builds semantic prototypes directly from the current tokens via a soft assignment. It then confirms them with a learnable query bank, matches tokens back to the confirmed prototypes with a learnable temperature, and propagates an explicit per-token confidence through ordering and aggregation. The resulting network, **DPRNet**, is trained on DIV2K and evaluated on five standard benchmarks at  $\times 2 / \times 3 / \times 4$ , where it maintains CATANet-level reconstruction quality with small gains on most benchmark cells and a comparable lightweight cost profile, while exposing an interpretable, confidence-aware routing mechanism.

## Contributions.

The main contributions are as follows.

- **C1 – Dynamic semantic prototypes.** We generate routing prototypes from the current tokens by a soft, assignment-weighted aggregation, removing the cross-batch EMA center buffer so that prototypes follow per-image semantics (Sec. 2.3, Eqs. (4)–(5)).
- **C2 – Decoupled prototype confirmation.** We separate prototype *generation* from *confirmation* by refining the content prototypes with a learnable query bank through a single gated cross-attention, improving the stability and interpretability of the routing slots (Eq. (7)).
- **C3 – Confidence-aware routing.** We explicitly retain the token assignment confidence and use it in three places: a composite ordering key that sorts high-confidence tokens first within each cluster, a confidence-modulated global gate, and a soft prototype fallback for uncertain tokens (Eqs. (10), (11), (12)).
- **C4 – Stabilized routing.** We introduce a learnable routing temperature together with a prototype usage-balancing regularizer to mitigate collapse onto a few prototypes and keep slot usage spread out (Eqs. (8), (13)).

Each design element is exposed as an independent switch that defaults to the full model. Turning these switches off gives the corresponding ablated variants, which allows the contributions to be isolated cleanly in the ablation study (Sec. 3.5).

## 2. Method 123

### 2.1. Overview 124

We build on the lightweight super-resolution (SR) backbone of CATANet and replace its token-routing component with a *Dynamic Prototype Routing* (DPR) module (Fig. 1). Given a low-resolution (LR) input  $I_{LR} \in \mathbb{R}^{3 \times H \times W}$ , a  $3 \times 3$  convolution extracts shallow features  $F_0 \in \mathbb{R}^{C \times H \times W}$  ( $C=40$ ). The deep feature extractor is a stack of  $L=8$  blocks; each block couples a content-routed global aggregation unit (the Token Aggregation Block, TAB) with a local self-attention unit (LRSA), followed by a  $3 \times 3$  fusion convolution and a residual connection:

$$F_l = F_{l-1} + \text{Conv}_{3 \times 3}(\text{LRSA}_l(\text{TAB}_l(F_{l-1}))), \quad l = 1, \dots, L. \quad (1) \quad 132$$

A global residual links the deep features to the network output, which is reconstructed by a pixel-shuffle upsampler and added to a bilinearly upsampled base image:

$$I_{SR} = \text{Up}(F_0 + F_L) + \text{Bilinear}_{\uparrow s}(I_{LR}), \quad (2) \quad 135$$

where  $s \in \{2, 3, 4\}$  is the scale factor. The proposed change is confined to TAB: we keep its external interface  $\mathbb{R}^{C \times H \times W} \rightarrow \mathbb{R}^{C \times H \times W}$  unchanged and replace the original means-based center iteration and hard-label sorting with the DPR pipeline (Sec. 2.3), while reusing the subsequent intra/inter-group aggregation operator (IASA, Sec. 2.4). This keeps the rest of the backbone untouched and isolates the effect of the routing mechanism.

Inside a TAB, the input feature map is flattened into a token sequence  $X \in \mathbb{R}^{B \times N \times C}$  with  $N=HW$ , layer-normalized, routed by DPR, aggregated by IASA, and fused by a  $1 \times 1$  convolution and a convolutional feed-forward network (ConvFFN), each with a residual connection. The data flow is

$$X \rightarrow \text{DPR} \rightarrow \text{IASA} \rightarrow \text{Conv}_{1 \times 1} \rightarrow \text{ConvFFN}. \quad (3) \quad 145$$

### 2.2. Motivation: limitations of center-based routing 146

The original TAB performs content-aware reordering by maintaining a set of routing centers via a  $k$ -means-style iteration with an exponential moving average (EMA) buffer, assigning each token to its nearest center by an  $\text{argmax}$  hard label, sorting tokens by that label, and aggregating the sorted sequence. This design leaves four weaknesses that DPR targets. (i) The centers are driven by a cross-batch historical buffer rather than strictly by the current input, so they adapt slowly to per-image semantics. (ii) The token-to-center relation is a hard label and discards assignment confidence. (iii) The ordering clusters by label only, with no intra-cluster purity constraint. (iv) Prototype generation and prototype confirmation are entangled, which weakens interpretability and makes the routing state harder to analyze. DPR addresses these issues by generating prototypes from the current tokens, decoupling generation from confirmation, and propagating an explicit assignment confidence through routing and aggregation.

### 2.3. Dynamic Prototype Routing 159

DPR maps a token sequence  $X \in \mathbb{R}^{B \times N \times C}$  to a content-reordered sequence,  $M$  dynamic prototypes, and routing metadata. The pipeline contains four stages (Fig. 2).

#### (1) Dynamic prototype generation. 162

We first embed tokens with a lightweight head  $E = \text{GELU}(\text{Linear}(\text{LN}(X))) \in \mathbb{R}^{B \times N \times C}$ , and predict a soft assignment of each token to  $M$  prototype slots:

$$A = \text{softmax}(EW_a) \in \mathbb{R}^{B \times N \times M}, \quad W_a \in \mathbb{R}^{C \times M}, \quad (4) \quad 165$$



where the softmax is taken over the  $M$  slots. Each prototype is the assignment-weighted average of the *current* tokens,

$$P_m^c = \frac{\sum_n A_{n,m} X_n}{\sum_n A_{n,m} + \epsilon}, \quad P^c = \text{normalize}(\text{LN}(P^c)), \quad (5)$$

with  $\ell_2$  normalization along the channel dimension. Unlike the EMA centers of the baseline,  $P^c$  depends only on the current input and carries no cross-batch state.

## (2) Prototype confirmation by query refinement.

To separate prototype *generation* from *confirmation*, we refine  $P^c$  with a learnable query bank  $Q \in \mathbb{R}^{M \times C}$  via a single cross-attention over the tokens. With query seed  $\tilde{P} = P^c + Q$ , queries from the seed, and keys/values from the embedded/raw tokens,

$$P^r = \text{softmax}\left(\frac{(\tilde{P}W_q)(EW_k)^\top}{\sqrt{d_r}}\right)(XW_v), \quad (6)$$

where  $d_r$  is the router dimension. The refinement is fused through a learnable gate  $\gamma = \sigma(g_r) \in (0, 1)$ ,

$$P = \text{normalize}(\text{LN}(P^c + \gamma P^r)). \quad (7)$$

Initializing  $g_r$  to a small value keeps the refinement a residual perturbation at the start of training, which helps keep prototype identities stable across samples. This stage can be disabled (ablation A1) by setting `use_prototype_query_refine=false`, which falls back to  $P = P^c$ .

## (3) Confidence-aware token-to-prototype matching.

We match tokens *back* to the confirmed prototypes in a shared  $\ell_2$ -normalized router space,  $\hat{e}_n = \text{normalize}(E_n W_t)$  and  $\hat{p}_m = \text{normalize}(P_m W_p)$ , and compute routing scores with a *learnable temperature*  $\tau$ :

$$S = \text{softmax}(\tau \hat{E} \hat{P}^\top) \in \mathbb{R}^{B \times N \times M}, \quad \tau = \min(e^\theta, \tau_{\max}), \quad (8)$$

where  $\theta$  is a learnable log-scale clamped by  $\tau_{\max}$  for stability. A larger  $\tau$  sharpens the assignment distribution and increases the separation of routing scores. For each token we read out the hard label and its confidence,

$$x_n^s = \max_m S_{n,m}, \quad b_n = \arg \max_m S_{n,m}, \quad (9)$$

and *explicitly retain* the confidence  $x_n^s$  (addressing the hard-label confidence discarded by TAB).

## (4) Confidence-aware ordering.

Instead of sorting only by the cluster label, we use a composite key that keeps the inter-cluster order while ordering high-confidence tokens first *within* each cluster:

$$k_n = b_n + 0.5(1 - x_n^s), \quad \pi = \text{argsort}_n(k_n). \quad (10)$$

Since  $x_n^s \in (0, 1)$ , the additive term lies in  $(0, 0.5)$  and never crosses an integer label boundary, so clusters stay contiguous while their members are ordered by confidence. The token sequence, labels, and confidences are gathered by  $\pi$  into  $X^\pi$ ,  $b^\pi$ ,  $x^{s,\pi}$ , and the inverse permutation  $\pi^{-1}$  is recorded for scatter-back. Setting `use_conf_sort=false` reduces Eq. (10) to a pure label sort  $k_n = b_n$  (ablation A2).

DPR outputs the reordered tokens  $X^\pi$ , the dynamic prototypes  $P$ , the full score map  $S$ , and the per-token confidence  $x^s$ , which are consumed by the aggregator below.

#### 2.4. Confidence-aware aggregation (IASA)

The aggregation operator is retained from the backbone but is made confidence-aware. On the reordered sequence  $X^\pi$  it computes (a) intra-group local attention over overlapping groups of size  $g$ , which captures fine local structure within each contiguous, semantically coherent segment; and (b) a global branch that attends to the dynamic prototypes  $P$  projected to keys/values. Let  $O_{\text{loc}}$  and  $O_{\text{glb}}$  denote the two branch outputs. They are combined through a learnable global gate  $\beta = \sigma(g_g)$  modulated by the per-token confidence:

$$O = O_{\text{loc}} + \beta \cdot x^{s,\pi} \cdot O_{\text{glb}}, \quad (11)$$

so that low-confidence tokens rely less on the global prototype summary. The output is scattered back to the original token order via  $\pi^{-1}$  and projected. Setting `use_iasa_score_gate=false` removes the confidence factor in Eq. (11) ( $O = O_{\text{loc}} + \beta O_{\text{glb}}$ ; ablation A2).

Soft prototype fallback.

To stabilize tokens whose hard assignment is uncertain, DPR adds a soft, confidence-weighted prototype context to the aggregated output. With a soft summary  $\bar{X} = SP$  projected by  $W_s$  and a learnable gate  $\alpha = \sigma(g_s)$ ,

$$Y = O + \alpha (1 - x^s) (\bar{X} W_s), \quad (12)$$

where  $(1 - x^s)$  emphasizes low-confidence tokens. The result  $Y$  passes through the  $1 \times 1$  convolution, residual connection, and ConvFFN. Setting `use_soft_fallback=false` reduces Eq. (12) to  $Y = O$  (ablation A2). Equations (10), (11) and (12) together form the confidence-aware routing path, and the three switches are independent and default to the full model.

#### 2.5. Prototype usage balancing

A learnable temperature can concentrate routing on a few prototypes. We therefore add a usage-entropy regularizer that encourages balanced slot use. Let  $u_m = \frac{1}{N} \sum_n S_{n,m}$  be the mean usage of prototype  $m$ ; its normalized entropy and the auxiliary loss are

$$\mathcal{H}(u) = -\frac{1}{\log M} \sum_m u_m \log u_m, \quad \mathcal{L}_{\text{bal}} = \lambda (1 - \mathcal{H}(u)), \quad (13)$$


which is largest, and therefore gives the smallest auxiliary loss, when usage is uniform. The total objective adds this term, summed over all blocks, to the pixel reconstruction loss:

$$\mathcal{L} = \mathcal{L}_{\text{pix}} + \sum_l \mathcal{L}_{\text{bal}}^{(l)}. \quad (14)$$

Setting the per-block weight  $\lambda=0$  disables balancing (ablation A3).

#### 2.6. Discussion of design

DPR keeps the backbone interface and parameter budget in the lightweight regime while changing only how tokens are organized before aggregation. The four stages map directly to the contributions: dynamic generation (Eqs. 4–5) for C1, generation/confirmation decoupling (Eq. 7) for C2, confidence-aware routing (Eqs. 8–10, 11–12) for C3, and the learnable temperature with usage balancing (Eqs. 8, 13) for C4. Because each design



**Figure 1.** Overall architecture of DPRNet, following the lightweight-SR backbone of CATANet. **(a)** A  $3 \times 3$  convolution extracts shallow features  $F_0$ , which pass through a stack of  $L=8$  DPRNet blocks; the deep features are added back through a global residual and reconstructed by a pixel-shuffle upsampler that is summed with a bilinearly upsampled base image. **(b)** Each block keeps the TAB–LRSA–fusion structure under a residual connection. **(c)** Inside a TAB, the only modified component is the routing module: DPR replaces center-based routing while preserving the input/output interface, and is expanded in Fig. 2.

element is exposed as an independent switch that defaults to the full model and has a defined disabled form, the contributions can be isolated cleanly in the ablation study (Sec. 3.5).

### 3. Experiments

#### 3.1. Datasets and implementation details

Datasets and metrics.


Following the standard lightweight-SR protocol, we train on the 800 training images of DIV2K [39] and evaluate on five benchmarks: Set5 [40], Set14 [41], BSD100 [42], Urban100 [43], and Manga109 [44]. Low-resolution inputs are produced by bicubic down-scaling at  $\times 2$ ,  $\times 3$ , and  $\times 4$ . We report PSNR and SSIM [45] on the Y channel of the YCbCr space, discarding a border of  $s$  pixels for scale  $s$ , which matches the convention used by the cited baselines.

Network configuration.

DPRNet keeps the CATANet backbone unchanged and replaces only the routing inside each TAB with the DPR module (Sec. 2). The deep feature extractor stacks  $L=8$  blocks with channel width  $C=40$ . The number of prototype slots per block is  $M = [16, 32, 64, 128, 16, 32, 64, 128]$ , and the intra-group attention uses a group size of 128. The learnable routing temperature is initialized at  $\tau=6$  (i.e. a log-scale parameter  $\theta = \log 6$ ) and clamped at  $\tau_{\max}=10$  (Eq. (8)); the per-block usage-balancing weights are  $\lambda = [5 \times 10^{-4}]_{\times 6}$  for the first six blocks and  $7 \times 10^{-4}$  for the last two (Eq. (13)). These yield 0.60–0.67 M parameters depending on scale (Table 2).

Training.

We train the  $\times 2$  model for 800k iterations and the  $\times 3$  and  $\times 4$  models for 250k iterations each. Training patches are  $128 \times 128$ ,  $192 \times 192$ , and  $256 \times 256$  HR crops for  $\times 2 / \times 3 / \times 4$



**Figure 2.** Dynamic Prototype Routing (DPR). **(a)** The router runs in a single pass: it embeds the tokens, forms content prototypes  $P^c$  from the current tokens, confirms them with a learnable query bank through a gated refinement ( $\gamma$ ), and matches the tokens back to the confirmed prototypes  $P$  with a bounded temperature  $\tau$ , yielding the ordered sequence  $X^\pi$ , the confidence  $x^s$ , and the label  $b$  (Eqs. (4)–(10)); the usage-balancing term acts on the score map  $S$  (Eq. (13)). **(b)** IASA consumes  $X^\pi$ ,  $P$ , and  $x^s$ : a local branch and a prototype global branch are combined under a confidence-modulated gate  $\beta$ , and a soft prototype fallback gated by  $\alpha$  stabilizes low-confidence tokens before the output is scattered back via  $\pi^{-1}$  (Eqs. (11)–(12)).

respectively, with a batch size of 16 and random horizontal flips and  $90^\circ$  rotations for augmentation. We optimize the L1 reconstruction loss together with the usage-balancing regularizer (Eq. (13)) using Adam ( $\beta_1=0.9$ ,  $\beta_2=0.99$ , no weight decay) at an initial learning rate of  $2 \times 10^{-4}$ , halved by a MultiStep schedule (milestones at  $\{300, 500, 650, 700, 750\}$ k for  $\times 2$  and  $\{125, 200, 225, 237.5\}$ k for  $\times 3 / \times 4$ ). The learnable routing temperature is trained with a  $0.1 \times$  learning-rate multiplier for stability. Validation is run every 2.5k iterations. Experiments use PyTorch with the BasicSR framework on a single NVIDIA RTX 4090 GPU.

### Efficiency measurement.

Parameters, Multi-Adds (MACs reported by thop), and inference latency are measured on a single NVIDIA RTX 4090 GPU for a fixed LR input of  $256 \times 256$  ( $\times 4$  output  $1024 \times 1024$ ), averaged over 100 runs after 20 warm-up runs. This input convention matches the one under which the CATANet paper reports Multi-Adds for itself and the lightweight Transformer baselines, making the figures directly comparable. Efficient-SR challenge reports likewise treat runtime, parameters, and FLOPs as joint deployment criteria rather than auxiliary statistics [22]. For comparison with FLOPs-based reports, Multi-Adds can be multiplied by two.

### 3.2. Comparison with state-of-the-art

We compare DPRNet with eight representative lightweight SR methods. The CNN group contains CARN [16], IMDN [17], RFDN [18], and RLFN [19], while the Transformer and routing group contains SwinIR-light [26], ELAN-light [27], SRFormer-light [28], and the direct predecessor CATANet [38]. Competitor numbers are cited from the original papers, as noted in Table 1. DPRNet is reported on all five benchmarks at each scale.

### Overall accuracy.

Table 1 shows that DPRNet and CATANet form the top tier at all three scales, with the clearest separation on Urban100 and Manga109 and smaller margins on Set5, Set14, and BSD100. For example, at  $\times 4$  DPRNet reaches 32.59 dB/0.9002 on Set5 and 27.76 dB/0.7434 on BSD100, compared with SRFormer-light at 32.51/0.8988 and 27.73/0.7422 under a comparable parameter budget. The pattern suggests that content-aware token aggregation is particularly useful in the lightweight regime when repeated or semantically related structures occur across distant positions.

### Comparison with CATANet.

Against CATANet—the strongest baseline and the method from which DPRNet differs only in the routing module—DPRNet ranks first or tied first on nearly every benchmark/scale cell, with CATANet usually second. The margins are small but directionally consistent: at  $\times 4$  DPRNet improves Urban100 from 26.87/0.8081 to 26.89/0.8089, Set5 from 32.58/0.8998 to 32.59/0.9002, Set14 from 28.90/0.7880 to 28.91/0.7883, and Manga109 PSNR from 31.31 to 31.32; at  $\times 2$  it reaches the best Set14 (34.06/0.9220), Urban100 PSNR (33.15), and Manga109 (39.38/0.9789). At  $\times 3$  it leads or ties CATANet on every benchmark (e.g. BSD100 29.29/0.8103 and Urban100 29.06/0.8689). The single metric on which CATANet remains ahead is Manga109 SSIM at  $\times 4$  (0.9183 vs. 0.9181). These gaps are within a few hundredths of a dB, so the result supports a conservative conclusion. Replacing the EMA-center routing of TAB with the dynamic prototype routing of DPR preserves CATANet-level reconstruction quality while changing the routing mechanism to the interpretable, confidence-aware formulation analyzed in Sec. 3.5 and Sec. 3.6.

### Efficiency.

Table 2 reports the  $\times 4$  efficiency subset for methods whose Multi-Adds are available under the same fixed  $256 \times 256$  LR-input convention. DPRNet has 660K parameters and 52.14 G Multi-Adds, keeping it within the lightweight bracket. It uses fewer Multi-Adds than the lightweight Transformer baselines SwinIR-light (60.3 G) and SRFormer-light (56.5 G) while delivering top-tier accuracy. Both its parameter count and Multi-Adds are modestly higher than CATANet's (660K/52.14 G vs. 535K/46.8 G at  $\times 4$ ), reflecting the learnable query bank and router projections of DPR. We therefore do not claim an efficiency advantage over CATANet. The per-scale DPRNet profile, including measured latency on our hardware, is given in Table 3. Cross-paper latency and compute values reported under incompatible conventions are excluded from Table 2 rather than mixed into the same numeric comparison.

### 3.3. Model complexity and efficiency

Table 2 compares model complexity at  $\times 4$  for the subset of methods whose Multi-Adds are available under the same input convention, and Table 3 gives DPRNet's per-scale profile. All DPRNet numbers are measured under the protocol of Sec. 3.1 with a fixed  $256 \times 256$  LR input ( $\times 4$  output  $1024 \times 1024$ ), reporting thop MACs and latency over 100 runs after 20 warm-ups. We use this  $256 \times 256$  input convention because it is the one under which the CATANet paper reports Multi-Adds for itself, SwinIR-light, and SRFormer-light; DPRNet's MACs are therefore directly comparable to those three cited figures. Methods whose published compute values use different conventions, or whose values are unavailable under this convention, are excluded from the efficiency subset rather than displayed as blank entries.

**Table 1.** Quantitative comparison with state-of-the-art lightweight SR methods on five benchmarks (PSNR(dB)/SSIM on the Y channel). **Best** and **second-best** results are highlighted. Competitor metrics are cited from the original papers (see notes). “-” denotes results not reported in the source.

| Method         | Params | Set5  |        | Set14 |        | BSD100 |        | Urban100 |        | Manga109 |        |
|----------------|--------|-------|--------|-------|--------|--------|--------|----------|--------|----------|--------|
|                |        | (K)   | PSNR   | SSIM  | PSNR   | SSIM   | PSNR   | SSIM     | PSNR   | SSIM     | PSNR   |
| Scale ×2       |        |       |        |       |        |        |        |          |        |          |        |
| CARN           | 1592   | 37.76 | 0.9590 | 33.52 | 0.9166 | 32.09  | 0.8978 | 31.92    | 0.9256 | 38.36    | 0.9765 |
| IMDN           | 694    | 38.00 | 0.9605 | 33.63 | 0.9177 | 32.19  | 0.8996 | 32.17    | 0.9283 | 38.88    | 0.9774 |
| RFDN           | 534    | 38.05 | 0.9606 | 33.68 | 0.9184 | 32.16  | 0.8994 | 32.12    | 0.9278 | 38.88    | 0.9773 |
| RLFN           | 527    | 38.07 | 0.9607 | 33.72 | 0.9187 | 32.22  | 0.9000 | 32.33    | 0.9299 | -        | -      |
| SwinIR-light   | 878    | 38.14 | 0.9611 | 33.86 | 0.9206 | 32.31  | 0.9012 | 32.76    | 0.9340 | 39.12    | 0.9783 |
| ELAN-light     | 582    | 38.17 | 0.9611 | 33.94 | 0.9207 | 32.30  | 0.9012 | 32.76    | 0.9340 | 39.11    | 0.9782 |
| SRFormer-light | 853    | 38.23 | 0.9613 | 33.94 | 0.9209 | 32.36  | 0.9019 | 32.91    | 0.9353 | 39.28    | 0.9785 |
| CATANet        | 477    | 38.28 | 0.9617 | 33.99 | 0.9217 | 32.37  | 0.9023 | 33.09    | 0.9372 | 39.37    | 0.9784 |
| DPRNet (ours)  | 602    | 38.29 | 0.9617 | 34.06 | 0.9220 | 32.38  | 0.9031 | 33.15    | 0.9372 | 39.38    | 0.9789 |
| Scale ×3       |        |       |        |       |        |        |        |          |        |          |        |
| CARN           | 1592   | 34.29 | 0.9255 | 30.29 | 0.8407 | 29.06  | 0.8034 | 28.06    | 0.8493 | 33.43    | 0.9427 |
| IMDN           | 703    | 34.36 | 0.9270 | 30.32 | 0.8417 | 29.09  | 0.8046 | 28.17    | 0.8519 | 33.61    | 0.9445 |
| RFDN           | 541    | 34.41 | 0.9273 | 30.34 | 0.8420 | 29.09  | 0.8050 | 28.21    | 0.8525 | 33.67    | 0.9449 |
| RLFN           | -      | -     | -      | -     | -      | -      | -      | -        | -      | -        | -      |
| SwinIR-light   | 886    | 34.62 | 0.9289 | 30.54 | 0.8463 | 29.20  | 0.8082 | 28.66    | 0.8624 | 33.98    | 0.9478 |
| ELAN-light     | 590    | 34.64 | 0.9288 | 30.55 | 0.8463 | 29.21  | 0.8081 | 28.69    | 0.8624 | 34.00    | 0.9478 |
| SRFormer-light | 861    | 34.67 | 0.9296 | 30.57 | 0.8469 | 29.26  | 0.8099 | 28.81    | 0.8655 | 34.19    | 0.9489 |
| CATANet        | 550    | 34.75 | 0.9300 | 30.67 | 0.8481 | 29.28  | 0.8101 | 29.04    | 0.8689 | 34.40    | 0.9499 |
| DPRNet (ours)  | 674    | 34.76 | 0.9301 | 30.67 | 0.8481 | 29.29  | 0.8103 | 29.06    | 0.8689 | 34.41    | 0.9499 |
| Scale ×4       |        |       |        |       |        |        |        |          |        |          |        |
| CARN           | 1592   | 32.13 | 0.8937 | 28.60 | 0.7806 | 27.58  | 0.7349 | 26.07    | 0.7837 | 30.42    | 0.9070 |
| IMDN           | 715    | 32.21 | 0.8948 | 28.58 | 0.7811 | 27.56  | 0.7353 | 26.04    | 0.7838 | 30.45    | 0.9075 |
| RFDN           | 550    | 32.24 | 0.8952 | 28.61 | 0.7819 | 27.57  | 0.7360 | 26.11    | 0.7858 | 30.58    | 0.9089 |
| RLFN           | 543    | 32.24 | 0.8952 | 28.62 | 0.7813 | 27.60  | 0.7364 | 26.17    | 0.7877 | -        | -      |
| SwinIR-light   | 897    | 32.44 | 0.8976 | 28.77 | 0.7858 | 27.69  | 0.7406 | 26.47    | 0.7980 | 30.92    | 0.9151 |
| ELAN-light     | 601    | 32.43 | 0.8975 | 28.78 | 0.7858 | 27.69  | 0.7406 | 26.54    | 0.7982 | 30.92    | 0.9150 |
| SRFormer-light | 873    | 32.51 | 0.8988 | 28.82 | 0.7872 | 27.73  | 0.7422 | 26.67    | 0.8032 | 31.17    | 0.9165 |
| CATANet        | 535    | 32.58 | 0.8998 | 28.90 | 0.7880 | 27.75  | 0.7427 | 26.87    | 0.8081 | 31.31    | 0.9183 |
| DPRNet (ours)  | 660    | 32.59 | 0.9002 | 28.91 | 0.7883 | 27.76  | 0.7434 | 26.89    | 0.8089 | 31.32    | 0.9181 |

Notes. Competitor numbers are cited verbatim: CARN/IMDN/RFDN/SwinIR-light/ELAN-light/CATANet from CATANet (CVPR 2025, Tab. 2); SRFormer-light from its paper (Tab. VI); RLFN from its paper (Tab. 1). RLFN does not report x3 or Manga109 in the original paper (“-”). Params for SwinIR-light and ELAN-light follow the convention of the CATANet table.

## Parameters.

DPRNet stays within the lightweight bracket at 660K parameters ( $\times 4$ ), comparable to Transformer baselines such as ELAN-light (601K) and well below SwinIR-light (897K), SRFormer-light (873K), and the CNN-heavy CARN (1592K). It is, however, larger than its direct predecessor CATANet (535K): the learnable query bank and the router/score projections of DPR add parameters relative to the EMA-center routing of CATANet. The comparison with CATANet is therefore framed as accuracy at a comparable or higher level, rather than as a parameter advantage (Urban100 26.89 vs. 26.87, with the full ablation analysis in Sec. 3.5). Fig. 3 plots this PSNR–parameter trade-off at  $\times 4$  on Urban100: DPRNet attains the highest PSNR among the compared lightweight methods while remaining within the same parameter bracket.

## Computation.

In terms of Multi-Adds, DPRNet requires 52.14 G at  $\times 4$  under the unified  $256 \times 256$  protocol. This is higher than CATANet’s 46.8 G, so the learnable query bank and the router/score projections of DPR do add inference cost relative to CATANet’s EMA-center routing. We therefore do not claim a compute advantage over CATANet; the relevant comparison is accuracy at a comparable or higher level (Urban100 26.89 vs. 26.87, full ablation in Sec. 3.5). The compute profile is more favorable relative to the Transformer baselines: at the same input size DPRNet uses markedly fewer Multi-Adds than both



**Table 2.** Comparable efficiency subset on  $\times 4$  SR. Multi-Adds are MACs under a unified fixed LR input of  $256 \times 256$  ( $\times 4$  output  $1024 \times 1024$ ); to compare against FLOPs-based reports, multiply by 2. Representative PSNR is on Urban100 ( $\times 4$ ).

| Method               | Params | MACs        | U100 PSNR    | Source         |
|----------------------|--------|-------------|--------------|----------------|
| SwinIR-light         | 897    | 60.3        | 26.47        | CATANet Tab. 6 |
| SRFormer-light       | 873    | 56.5        | 26.67        | CATANet Tab. 6 |
| CATANet              | 535    | <b>46.8</b> | <u>26.87</u> | CATANet Tab. 6 |
| <b>DPRNet (ours)</b> | 660    | 52.14       | <b>26.89</b> | measured       |

*Notes.* All Multi-Adds are MACs under a common  $256 \times 256$  LR input ( $\times 4$  output  $1024 \times 1024$ ): DPRNet is measured under this protocol, and the CATANet, SwinIR-light and SRFormer-light values are taken from the CATANet paper (CVPR'25, Tab. 6), which uses the identical input convention, so the four entries are directly comparable. DPRNet needs more Multi-Adds than its direct predecessor CATANet (52.14 vs. 46.8 G) but markedly fewer than the Transformer baselines SwinIR-light (60.3 G) and SRFormer-light (56.5 G), while attaining the best Urban100 PSNR in this subset. Methods whose MACs are not reported under this shared input convention are omitted from the efficiency subset rather than filled from incompatible sources; the complete accuracy and parameter comparison remains in Table 1. Cross-paper latency is also omitted because runtime is hardware- and implementation-dependent.

**Table 3.** DPRNet per-scale efficiency profile (measured). All scales use a fixed LR input of  $256 \times 256$ ; Multi-Adds are MACs (thop); latency averaged over 100 runs after 20 warm-ups on a single CUDA GPU.

| Scale      | LR Input         | Params (K) | Multi-Adds (G) | Latency (ms)      |
|------------|------------------|------------|----------------|-------------------|
| $\times 2$ | $256 \times 256$ | 601.95     | 36.19          | $183.57 \pm 7.70$ |
| $\times 3$ | $256 \times 256$ | 674.15     | 41.26          | $185.88 \pm 6.33$ |
| $\times 4$ | $256 \times 256$ | 659.71     | 52.14          | $198.16 \pm 9.43$ |

SwinIR-light (60.3 G) and SRFormer-light (56.5 G), while reaching the best Urban100 PSNR in Table 2. DPRNet thus remains firmly inside the lightweight bracket, trading a modest cost increase over its direct predecessor for an interpretable, confidence-aware routing mechanism and a clear cost reduction relative to attention-heavy Transformer designs at higher accuracy.

Per-scale behavior.


Table 3 reports the profile at a fixed  $256 \times 256$  LR input. Because the input size is held constant, the backbone cost is shared across scales and the Multi-Adds are dominated by the upsampling tail, so they *grow* with the scale factor—from 36.19 G ( $\times 2$ ) to 41.26 G ( $\times 3$ ) and 52.14 G ( $\times 4$ ). Latency follows the same order (183.57, 185.88, 198.16 ms). The parameter count is essentially scale-independent ( $\approx 0.60$ – $0.67$  M), as only the input projection and the upsampling tail change with scale.

Scope of the efficiency comparison.

The main accuracy table (Table 1) remains the full comparison against all eight baselines. The efficiency table is deliberately narrower: it reports only common-protocol MACs plus DPRNet's measured per-scale runtime. We do not mix latency values across different devices or implementation stacks, and we do not import compute numbers reported under incompatible input-size or FLOPs-counting conventions. This makes Table 2 a directly comparable efficiency subset rather than an all-method table with non-comparable entries.

### 3.4. Qualitative comparison

Beyond aggregate PSNR/SSIM, Fig. 4 compares reconstructions at  $\times 4$  on three texture-dense hard samples selected by ground-truth high-frequency energy: two Urban100 images dominated by repetitive building grids (img\_092, img\_024) and one Manga109 page of




**Figure 3.** PSNR vs. parameter count at  $\times 4$  on Urban100 (numbers from Table 2; competitor params cited from the original papers). DPRNet (\*) attains the highest PSNR among the compared lightweight methods while staying within the same parameter bracket (660K), trading a small parameter increase over its predecessor CATANet for the best reconstruction quality.

high-contrast line art (That's Izumiko). These are exactly the regimes in which content-based routing is expected to help most, since distant tokens that share the same structure can be aggregated together rather than only within a local window.

On the repetitive facades of Urban100, the bicubic baseline collapses adjacent window edges into blurred, aliased bands, and the lightweight CNN IMDN only partially recovers them. The transformer-based methods (SwinIR-light, SRFormer-light, CATANet) progressively sharpen the grid structure, and DPRNet reconstructs the window spacing with clean, well-separated boundaries on par with the strongest content-routing baseline, CATANet. On the Manga109 line art, where thin strokes are easily broken by upsampling, DPRNet keeps the contours continuous and avoids obvious ringing, matching the visual quality of the best transformer baselines. These observations are consistent with the routing diagnostics of Sec. 3.6: confidence-aware aggregation tracks structured, high-frequency regions that local-window methods resolve less faithfully. We stress that Fig. 4 is a qualitative comparison; the corresponding quantitative results for all baselines are reported in Table 1, where DPRNet remains in the top tier at a comparable parameter budget.

### 3.5. Ablation study

We ablate the four design elements of DPR using the independent switches exposed in Sec. 2. All ablations are conducted at scale  $\times 4$ , following CATANet [38] whose ablations are likewise performed at  $\times 4$ , the setting in which content routing over high-frequency texture is most stressed. Each variant is trained for 250k iterations on DIV2K under the same schedule as the full model and repeated over three seeds {3407, 42, 1234}. We report the mean  $\pm$  std across the three seeds at the validation-best checkpoint, selected by the Set5+Urban100 mean on the Y channel with matched crop\_border. Besides PSNR/SSIM, we log routing diagnostics per TAB block—the normalized prototype-usage entropy  $\mathcal{H}(u)$  of Eq. (13), the active-prototype usage, and the learnable temperature—to inspect the routing behavior directly. Because every variant is trained independently under the same



**Figure 4.** Qualitative comparison at  $\times 4$  on texture-dense hard samples. The top reference strip shows the ground-truth images with the crop regions; the rows below magnify those regions for Bicubic, IMDN, SwinIR-light, SRFormer-light, CATANet, DPRNet (ours), and the ground truth. DPRNet restores repetitive building grids and high-contrast line art with clear structure, visually comparable to the strongest content-routing baseline on these challenging crops. The samples are chosen by high-frequency content rather than by per-crop DPRNet advantage, and quantitative PSNR/SSIM are reported in Table 1.

protocol, the differences below reflect the contribution of the switched element itself. The three-seed protocol also lets us compare the mean effect with its run-to-run stability.

C1: dynamic prototypes vs. history centers.

Table 4 directly compares DPR's dynamic, per-image prototype generation (the full model) against a cross-batch history center with an EMA update (*EMA-center*), a re-implementation of the center-based mechanism of the baseline that keeps the identical backbone, interface, and everything downstream of the prototype generator. Under this matched protocol, dynamic prototypes give the better reconstruction on both benchmarks (Set5 32.57 vs. 32.51 dB, Urban100 26.87 vs. 26.76 dB) and reduce seed-to-seed variance. This provides a direct controlled confirmation of C1: forming prototypes from the current tokens, rather than accumulating them across batches, is what improves the routing—not merely a difference in capacity or training budget, which are held fixed here.

A1: decoupled prototype confirmation.

Table 5 compares DPR with and without the query-refinement stage (Eq. (7)). Enabling refine improves the result on both benchmarks (Set5 32.57 vs. 32.53 dB; Urban100 26.87 vs. 26.82 dB), and because the gain ( $\sim 0.04$ – $0.05$  dB) exceeds the seed-to-seed standard deviation ( $\sim 0.02$ – $0.04$  dB), it is attributable to the refine switch itself rather than to run-to-run noise. The refinement is gated by a small-initialized scalar (Eq. (7)), so it acts as a

residual confirmation step that keeps the prototype slots stable while decoupling generation from confirmation (C2); we therefore retain it by default.

A2: confidence-aware routing.

Table 6 adds the three confidence-aware components cumulatively: confidence-aware ordering (Eq. (10)), the IASA score gate (Eq. (11)), and the soft prototype fallback (Eq. (12)). The full configuration (v4) is clearly the best of the four on both Set5 and Urban100, and by more than the seed standard deviation. The intermediate steps, however, are not monotone—adding the ordering and gate alone slightly lowers PSNR before the complete set recovers and surpasses the hard-routing baseline. We therefore read the three switches as mutually compatible and most effective as the *complete* confidence-aware path, rather than as three individually additive accuracy gains. We keep the components because they make the routing path interpretable and confidence-driven: the retained per-token confidence  $x^s$  modulates ordering, the global gate, and the fallback (C3), which we visualize through the  $x^s$  distributions in Sec. 3.6.

A3: prototype usage balancing.

Table 7 isolates the usage-balancing regularizer (Eq. (13)). Enabling it raises the block-averaged normalized usage entropy from 0.623 to 0.686 on Urban100, and the increase is *consistent across all eight TAB blocks* ( $b_0$ – $b_7$ ; Table 8), i.e. the regularizer drives a more balanced use of prototype slots and counteracts the collapse that the learnable temperature would otherwise encourage (C4). Reconstruction quality is also slightly higher with the regularizer enabled (Set5 32.57 vs. 32.55 dB; Urban100 26.87 vs. 26.85 dB), so the balancing does not trade accuracy for a flatter usage distribution. With three seeds available, we can also assess the second part of C4: enabling the balance loss *reduces* the seed-to-seed standard deviation on both benchmarks ( $0.05 \rightarrow 0.02$  dB) and on the usage entropy itself ( $0.008 \rightarrow 0.003$ ). The regularizer thus makes both the accuracy and the routing behavior more stable across seeds, supporting the variance-reduction sub-claim of C4 rather than only the balancing mechanism.

Summary.

Across three seeds, the ablations show that DPR’s design elements are mutually compatible: the full model is the best variant in every group, and each component either preserves or improves reconstruction quality. Beyond the mean effects, the study substantiates two mechanism claims: dynamic prototypes outperform an EMA history center under a matched protocol (C1), and the usage-balancing regularizer both increases prototype-usage entropy block-consistently and lowers seed-to-seed variance (C4).

### 3.6. Routing interpretability analysis

To inspect DPR’s routing behavior independently of its effect on PSNR, we analyze the diagnostics logged during validation at  $\times 4$  (Urban100, full model). For each of the  $L=8$  TAB blocks we record the per-token routing confidence  $x^s = \max_m S_{n,m}$  (Eq. (9)), the number of *active* prototypes (slots receiving at least one token), the normalized usage entropy  $\mathcal{H}(u)$  (Eq. (13)), and the learned temperature  $\tau$  (Eq. (8)). The blocks use  $M = [16, 32, 64, 128, 16, 32, 64, 128]$  prototypes respectively.

Balanced prototype usage (C4).

Fig. 6 reports the normalized usage entropy per block with and without the balancing regularizer. With balancing enabled (the full model) the entropy is higher in *all eight* blocks, raising the block-mean from 0.623 to 0.686; the effect is largest in the high-frequency-oriented block  $b_4$  ( $0.709 \rightarrow 0.847$ ). Consistently, the average number of active prototypes

**Table 4.** Ablation C1 ( $\times 4$ ): dynamic per-image prototypes (DPR, = full model) vs. a cross-batch history center with EMA update (EMA-center), both trained for 250k on DIV2K over three seeds. Values are mean $\pm$ std PSNR(dB)/SSIM (Y channel) on Set5 and Urban100. Under a matched protocol that differs *only* in the prototype generator, dynamic prototypes give the better reconstruction on both benchmarks with lower seed variance.

| Benchmark | Prototype generator  | PSNR                               | SSIM                                  |
|-----------|----------------------|------------------------------------|---------------------------------------|
| Set5      | EMA-center           | 32.51 $\pm$ 0.04                   | 0.8990 $\pm$ 0.0004                   |
| Set5      | DPR (dynamic, =full) | <b>32.57 <math>\pm</math> 0.02</b> | <b>0.9000 <math>\pm</math> 0.0002</b> |
| Urban100  | EMA-center           | 26.76 $\pm$ 0.06                   | 0.8052 $\pm$ 0.0009                   |
| Urban100  | DPR (dynamic, =full) | <b>26.87 <math>\pm</math> 0.02</b> | <b>0.8085 <math>\pm</math> 0.0003</b> |

*Notes.* Both arms share the identical CATANet backbone and DPR interface and differ only in how routing prototypes are formed: EMA-center maintains a cross-batch  $k$ -means-style center with an exponential-moving-average buffer and a hard  $\arg\max$  label (EMACenterRouter), whereas DPR forms prototypes from the current tokens by a soft, assignment-weighted aggregation. Each variant is trained for 250k iterations on DIV2K with seeds {3407, 42, 1234}; we report the mean $\pm$ std over the three seeds. Dynamic prototypes improve PSNR on both benchmarks (Set5 +0.06 dB, Urban100 +0.11 dB) and exhibit lower seed-to-seed variance, giving a direct controlled confirmation of C1.

**Table 5.** Ablation A1 ( $\times 4$ ): prototype query refine. PSNR(dB) mean $\pm$ std (Y channel) over three seeds. Enabling refine improves reconstruction on both benchmarks by more than the seed variance, confirming that decoupling prototype generation from confirmation (C2) is quality-positive; we retain it by default.

| Variant             | Set5 PSNR                          | Urban100 PSNR                      |
|---------------------|------------------------------------|------------------------------------|
| w/o refine          | 32.53 $\pm$ 0.03                   | 26.82 $\pm$ 0.04                   |
| w/ refine (default) | <b>32.57 <math>\pm</math> 0.02</b> | <b>26.87 <math>\pm</math> 0.02</b> |

*Notes.* Each variant is trained for 250k iterations on DIV2K with seeds {3407, 42, 1234} and reported as mean $\pm$ std at the validation-best checkpoint. Enabling refine yields the better PSNR on both Set5 (+0.04 dB) and Urban100 (+0.05 dB); because the improvement exceeds the seed-to-seed standard deviation, it is attributable to the query-refine switch itself. The comparison of dynamic prototypes against a history center with EMA update (C1) is reported separately in Table 4.

increases from 47.5 to 52.3 across blocks, meaning that fewer slots are left unused. Because a learnable temperature can otherwise sharpen the assignment distribution and concentrate tokens on a handful of prototypes, these two measurements together show that the regularizer keeps prototype usage spread out and prevents slot collapse, matching the stabilizing role claimed for C4.

Bounded routing temperature (C4).

The learnable temperature is clamped by  $\tau_{\max}=10$  for stability (Eq. (8)). In the trained full model it settles in a moderate range across blocks (per-block  $\tau$  from 5.65 to 6.51, block-mean 6.06), staying well below the clamp rather than saturating it. The temperature therefore sharpens the routing scores enough to separate prototypes while the usage regularizer keeps the assignment from collapsing—the two C4 components act as intended without hitting their safety bound.

Confidence-aware routing signal (C3).

The per-token confidence  $x^s$  that drives the ordering key (Eq. (10)), the IASA gate (Eq. (11)) and the soft fallback (Eq. (12)) is a non-degenerate signal in every block. Fig. 5 shows its distribution accumulated over *all* test images of Urban100 and Manga109 with the final model: in each block the probability mass lies clearly to the right of the uniform-assignment floor  $1/M$  (for example, the Urban100 block means are 0.109 vs. 0.0625 for  $b_0$  and 0.012 vs. 0.0078 for  $b_3$ ) yet stays far from 1. The softmax assignment is therefore informative but soft rather than collapsed onto a single prototype. As  $M$  grows, its mode shifts toward the lower floor. Reshaping the original-order labels  $b_n = \arg \max_m S_{n,m}$



**Table 6.** Ablation A2 ( $\times 4$ ): incremental confidence-aware routing components. PSNR(dB) mean $\pm$ std (Y channel) over three seeds. The full configuration (v4) is clearly the best on both benchmarks, but the intermediate steps are not monotone; the components are mutually compatible and best as a complete confidence-aware path rather than individually accuracy-additive.

| Variant             | Switches |      |    | Set5                               | Urban100                           |
|---------------------|----------|------|----|------------------------------------|------------------------------------|
|                     | sort     | gate | fb | PSNR                               | PSNR                               |
| v1 (hard sort)      | ✗        | ✗    | ✗  | 32.54 $\pm$ 0.03                   | 26.81 $\pm$ 0.02                   |
| v2 (+conf. sort)    | ✓        | ✗    | ✗  | 32.52 $\pm$ 0.03                   | 26.78 $\pm$ 0.03                   |
| v3 (+score gate)    | ✓        | ✓    | ✗  | 32.53 $\pm$ 0.03                   | 26.80 $\pm$ 0.04                   |
| v4 (full, +soft fb) | ✓        | ✓    | ✓  | <b>32.57 <math>\pm</math> 0.02</b> | <b>26.87 <math>\pm</math> 0.02</b> |

Notes. “sort”=confidence-aware sort key, “gate”=IASA score gating, “fb”=soft fallback. Each variant is trained for 250k iterations on DIV2K with seeds {3407, 42, 1234} and reported as mean $\pm$ std at the validation-best checkpoint.

The full configuration (v4) is the best of the four on both benchmarks and by more than the seed standard deviation, but the intermediate steps (v2/v3) are not monotone; the components are therefore reported as mutually compatible and best as the complete confidence-aware path, rather than as a monotone per-switch accuracy gain.

**Table 7.** Ablation A3 ( $\times 4$ ): prototype balance loss. Reconstruction PSNR(dB) and block-averaged normalized prototype-usage entropy on Urban100 (higher = more balanced), each reported as mean $\pm$ std over three seeds. The balance loss raises usage entropy while also *reducing seed-to-seed variance* on both benchmarks and on the usage entropy itself, supporting the stabilizing role claimed for C4.

| Variant                | Usage entropy ( $\uparrow$ )        | Set5 PSNR                          | Urban100 PSNR                      |
|------------------------|-------------------------------------|------------------------------------|------------------------------------|
| w/o balance loss       | 0.623 $\pm$ 0.008                   | 32.55 $\pm$ 0.05                   | 26.85 $\pm$ 0.05                   |
| w/ balance loss (def.) | <b>0.686 <math>\pm</math> 0.003</b> | <b>32.57 <math>\pm</math> 0.02</b> | <b>26.87 <math>\pm</math> 0.02</b> |

**Table 8.** Per-block normalized prototype-usage entropy (Urban100, representative seed 3407). The balance loss increases usage entropy in every TAB block ( $b_0$ – $b_7$ ).

| Variant     | $b_0$ | $b_1$ | $b_2$ | $b_3$ | $b_4$ | $b_5$ | $b_6$ | $b_7$ |
|-------------|-------|-------|-------|-------|-------|-------|-------|-------|
| w/o balance | 0.648 | 0.566 | 0.695 | 0.627 | 0.709 | 0.576 | 0.628 | 0.533 |
| w/ balance  | 0.718 | 0.585 | 0.723 | 0.681 | 0.847 | 0.635 | 0.685 | 0.619 |


Notes. Each variant is trained for 250k iterations on DIV2K with seeds {3407, 42, 1234}; reconstruction and the block-mean usage entropy are reported as mean $\pm$ std over the three seeds, and the per-block entropies are shown for the representative seed 3407. The balance loss consistently increases prototype-usage entropy across all blocks (more balanced usage), supporting C4’s balancing mechanism. With three seeds available, enabling the balance loss additionally lowers the seed-to-seed standard deviation on both benchmarks (0.05  $\rightarrow$  0.02 dB) and on the usage entropy itself (0.008  $\rightarrow$  0.003), so C4’s variance-reduction sub-claim is now supported directly.

to  $H \times W$  yields the routing cluster maps in Fig. 7: the partitions track image content (repetitive building facades, high-contrast strokes) instead of a fixed spatial grid. In this visualization the important evidence is boundary continuity rather than the exact hue: early blocks form coarse content groups, while deeper larger- $M$  blocks split the repeated structures into finer partitions. Both views are produced by a single forward pass with the cached routing state (last\_routing\_map, last\_x\_scores) and together confirm that the confidence used by the C3 components is meaningful.

### Summary.

The logged diagnostics provide mechanism-level evidence for the routing design: the balancing regularizer produces a block-consistent increase in usage entropy and active-prototype count (C4), the learnable temperature stays in a stable, unsaturated range (C4), and the retained confidence signal is informative across all blocks (C3). These observations characterize the behavior of DPR independently of the reconstruction metrics discussed in Sec. 3.5.





**Figure 5.** Distribution of the per-token routing confidence  $x^s = \max_m S_{n,m}$  (Eq. (9)) at  $\times 4$ , accumulated over *all* test images of (a) Urban100 and (b) Manga109 with the final model. Curves are shown for three representative TAB blocks with increasing prototype count  $M$ ; dotted vertical lines mark the uniform-assignment floor  $1/M$ . In every block the distribution sits clearly to the right of its floor (block-mean  $\approx 1.5\text{--}1.9 \times 1/M$ ) yet stays far from 1, confirming that the softmax assignment is *informative but soft*: tokens commit to a prototype more than chance while retaining the graded confidence that the C3 components (ordering key, IASA gate, soft fallback) rely on. As  $M$  grows the mode shifts toward the (lower) floor, consistent with confidence being spread over more slots.


## 4. Discussion

Content routing vs. spatial locality.

The benchmark results in Sec. 3.2 indicate where content-aware routing helps most. On Urban100 and Manga109—images dominated by repetitive structures and long-range self-similarity—DPRNet and CATANet, the two routing methods, separate more noticeably from the window- and distillation-based baselines, whereas on the more stochastic textures of BSD100 the margin between all top-tier methods narrows. This is consistent with the premise of content routing and with earlier SR work that exploits internal self-similarity or cross-scale non-local correspondences [2,35,43]: grouping tokens by semantics is most useful when semantically related content is spatially scattered. It is less useful when local statistics already capture the signal. DPR does not discard spatial locality; it complements it, since each TAB is paired with a local self-attention unit (Sec. 2.1); the routing module supplies the long-range, content-based interaction that a fixed spatial window cannot.

Small margins and evidence strength.

The numerical margins over CATANet are deliberately framed as small. The value of DPR is therefore not a claim of large PSNR gain over its direct predecessor, but a combination of three evidence streams: a consistent top-tier position across five benchmarks and three scales (Sec. 3.2), a preserved lightweight profile relative to attention-heavy Transformer baselines (Sec. 3.3), and explicit routing diagnostics showing meaningful confidence scores and more balanced prototype usage (Sec. 3.6). This framing matters because modern lightweight SR methods are often separated by only a few hundredths of a dB under standard PSNR/SSIM evaluation; without mechanism-level analysis, such small differences would be weak evidence on their own.



**Figure 6.** Normalized prototype-usage entropy  $\mathcal{H}(u)$  (Eq. (13); higher = more balanced) per TAB block ( $b_0$ – $b_7$ ) at  $\times 4$  on Urban100 (full model, representative seed 3407). The usage-balancing regularizer increases the entropy in *all eight* blocks (block-mean  $0.623 \rightarrow 0.686$ ), confirming that it counteracts the collapse onto a few prototypes encouraged by the learnable temperature (C4).

The risk of over-strong routing.


A confidence-aware router can, in principle, sharpen its assignments until tokens collapse onto a few prototypes, which would defeat the purpose of routing and destabilize training. Two components of DPR are explicitly designed against this failure mode: the routing temperature is clamped (Eq. (8)) and the usage-balancing regularizer (Eq. (13)) penalizes peaky usage distributions. The diagnostics in Sec. 3.6 show both mechanisms operating as intended—the learned temperature settles in a moderate, unsaturated range (block-mean 6.06, below the clamp of 10) and the balancing regularizer raises the prototype-usage entropy in all eight blocks ( $0.623 \rightarrow 0.686$ ) while increasing the number of active prototypes. Thus, the benefits of confidence-aware routing are obtained without letting the router become degenerate.

Prototype count and the cost trade-off.

DPR uses a per-block prototype budget  $M = [16, 32, 64, 128, 16, 32, 64, 128]$ , growing within each stage. More prototypes give the router a finer content vocabulary but enlarge the query bank and the matching cost, while too few prototypes force unrelated tokens to share a slot. We keep the prototype schedule aligned with the CATANet backbone so that the comparison isolates the routing mechanism rather than a capacity change; this is also why DPRNet carries a modestly higher parameter count and Multi-Adds than CATANet (660K/52.14 G vs. 535K/46.8 G at  $\times 4$  under a common  $256 \times 256$  input), as the added cost lives in the query bank and projections of DPR. Even so, DPRNet stays within the lightweight bracket and below the Transformer baselines SwinIR-light (60.3 G) and SRFormer-light (56.5 G) while reaching higher Urban100 PSNR (Sec. 3.3). A systematic sweep of the prototype schedule is a natural extension but is orthogonal to the routing formulation studied here.

Limitations.

Several limitations bound the present study and frame our claims conservatively. The ablations (Sec. 3.5) train every variant over three seeds, which supports single-switch attribution and lets us report both the mean effect and its seed-to-seed variance; the margins over the strongest configuration nonetheless remain small, so we frame the design elements as mutually compatible and at least quality-preserving rather than as sources of large individual gains. The direct comparison against a re-implemented EMA center branch (C1) is conducted only at  $\times 4$  on Set5 and Urban100, the setting in which content routing is most stressed, rather than across all five benchmarks. The efficiency comparison is



**Figure 7.** Routing cluster maps at  $\times 4$ . Each LR pixel is colored by its assigned prototype index  $b_n = \arg \max_m S_{n,m}$  (Eq. (9)) and reshaped from the original token order to  $H \times W$ . Read each row by comparing the cluster boundaries with the reference LR image: the early block forms coarse content groups, while the deeper blocks split repeated facade elements and high-contrast strokes into finer partitions. Colors encode prototype identity within a block and are not comparable across blocks, so the meaningful cue is region continuity rather than a specific hue. This visualizes the same routing signal whose confidence  $x^s$  is quantified in Fig. 5.

intentionally limited to common-protocol MACs and our measured DPRNet runtime: we do not claim cross-method latency because the competing models were not remeasured on the same hardware and implementation stack (Sec. 3.3). The models are trained on DIV2K only; larger training sets (e.g. DF2K) and real-world or arbitrary-scale degradations are not explored here and may further shift the trade-off.

## 5. Conclusion

We presented Dynamic Prototype Routing (DPR), a drop-in replacement for the center-based Token Aggregation Block of CATANet that keeps the backbone and tensor interface unchanged while remaining in a comparable lightweight operating regime. DPR generates semantic prototypes directly from the current tokens, decouples their confirmation through a learnable query bank, matches tokens back with a learnable temperature, and propagates explicit per-token confidence through ordering, gating, and a soft fallback. A usage-balancing regularizer keeps prototype use spread out. On five standard benchmarks at  $\times 2 / \times 3 / \times 4$ , the resulting DPRNet maintains CATANet-level reconstruction quality with small gains on most benchmark cells and a comparable lightweight cost profile. Across three seeds, the ablations attribute the effect of each design element and show that dynamic prototypes outperform a re-implemented EMA history center. The routing diagnostics further provide mechanism-level evidence that the temperature remains unsaturated and

that prototype usage becomes measurably more balanced—and more stable across seeds—in all blocks. Because DPR preserves the external interface of the routing block, it is not tied to this backbone and can serve as a confidence-aware content-routing module for other token-aggregation architectures. Future work includes remeasuring all comparison models for cross-method runtime under a single implementation protocol, and extending DPR to larger training sets and real-world or arbitrary-scale degradations.

**Author Contributions:** Conceptualization, methodology, software, validation, and writing—original draft preparation: Author Name. All authors have read and agreed to the published version of the manuscript.

**Funding:** This research received no external funding.

**Data Availability Statement:** The DIV2K training set and the Set5, Set14, BSD100, Urban100, and Manga109 benchmarks used in this study are publicly available. Code and trained models will be released upon publication.

**Conflicts of Interest:** The authors declare no conflicts of interest.

## References

1. Yang, J.; Wright, J.; Huang, T.S.; Ma, Y. Image Super-Resolution via Sparse Representation. *IEEE Transactions on Image Processing* **2010**, *19*, 2861–2873. <https://doi.org/10.1109/TIP.2010.2050625>.
2. Glasner, D.; Bagon, S.; Irani, M. Super-Resolution from a Single Image. In Proceedings of the IEEE International Conference on Computer Vision (ICCV), 2009, pp. 349–356. <https://doi.org/10.1109/ICCV.2009.5459271>.
3. Dong, C.; Loy, C.C.; He, K.; Tang, X. Learning a Deep Convolutional Network for Image Super-Resolution. In Proceedings of the European Conference on Computer Vision (ECCV), 2014, pp. 184–199. [https://doi.org/10.1007/978-3-319-10593-2\\_13](https://doi.org/10.1007/978-3-319-10593-2_13).
4. Kim, J.; Lee, J.K.; Lee, K.M. Accurate Image Super-Resolution Using Very Deep Convolutional Networks. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016, pp. 1646–1654. arXiv:1511.04587, <https://doi.org/10.1109/CVPR.2016.182>.
5. Kim, J.; Lee, J.K.; Lee, K.M. Deeply-Recursive Convolutional Network for Image Super-Resolution. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016, pp. 1637–1645. arXiv:1511.04491, <https://doi.org/10.1109/CVPR.2016.181>.
6. Lai, W.S.; Huang, J.B.; Ahuja, N.; Yang, M.H. Deep Laplacian Pyramid Networks for Fast and Accurate Super-Resolution. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2017, pp. 5835–5843. <https://doi.org/10.1109/CVPR.2017.618>.
7. Haris, M.; Shakhnarovich, G.; Ukita, N. Deep Back-Projection Networks for Super-Resolution. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2018, pp. 1664–1673. arXiv:1803.02735, <https://doi.org/10.1109/CVPR.2018.00179>.
8. Zhang, Y.; Tian, Y.; Kong, Y.; Zhong, B.; Fu, Y. Residual Dense Network for Image Super-Resolution. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2018, pp. 2472–2481. arXiv:1802.08797, <https://doi.org/10.1109/CVPR.2018.00262>.
9. Zhang, Y.; Li, K.; Li, K.; Wang, L.; Zhong, B.; Fu, Y. Image Super-Resolution Using Very Deep Residual Channel Attention Networks. In Proceedings of the European Conference on Computer Vision (ECCV), 2018, pp. 294–310. arXiv:1807.02758, [https://doi.org/10.1007/978-3-030-01234-2\\_18](https://doi.org/10.1007/978-3-030-01234-2_18).
10. Lim, B.; Son, S.; Kim, H.; Nah, S.; Lee, K.M. Enhanced Deep Residual Networks for Single Image Super-Resolution. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), 2017, pp. 1132–1140. arXiv:1707.02921, <https://doi.org/10.1109/CVPRW.2017.151>.
11. Wang, Z.; Chen, J.; Hoi, S.C.H. Deep Learning Algorithms for Single Image Super-Resolution: A Systematic Review. *Electronics* **2021**, *10*, 867. <https://doi.org/10.3390/electronics10070867>.
12. Dong, C.; Loy, C.C.; Tang, X. Accelerating the Super-Resolution Convolutional Neural Network. In Proceedings of the European Conference on Computer Vision (ECCV), 2016, pp. 391–407. arXiv:1608.00367, [https://doi.org/10.1007/978-3-319-46475-6\\_25](https://doi.org/10.1007/978-3-319-46475-6_25).



13. Shi, W.; Caballero, J.; Huszar, F.; Totz, J.; Aitken, A.P.; Bishop, R.; Rueckert, D.; Wang, Z. Real-Time Single Image and Video Super-Resolution Using an Efficient Sub-Pixel Convolutional Neural Network. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016, pp. 1874–1883. arXiv:1609.05158, <https://doi.org/10.1109/CVPR.2016.207>.
14. Tai, Y.; Yang, J.; Liu, X. Image Super-Resolution via Deep Recursive Residual Network. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2017, pp. 2790–2798. <https://doi.org/10.1109/CVPR.2017.298>.
15. Tai, Y.; Yang, J.; Liu, X.; Xu, C. MemNet: A Persistent Memory Network for Image Restoration. In Proceedings of the IEEE International Conference on Computer Vision (ICCV), 2017, pp. 4549–4557. arXiv:1708.02209, <https://doi.org/10.1109/ICCV.2017.486>.
16. Ahn, N.; Kang, B.; Sohn, K.A. Fast, Accurate, and Lightweight Super-Resolution with Cascading Residual Network. In Proceedings of the European Conference on Computer Vision (ECCV), 2018, pp. 256–272. arXiv:1803.08664, [https://doi.org/10.1007/978-3-030-01249-6\\_16](https://doi.org/10.1007/978-3-030-01249-6_16).
17. Hui, Z.; Gao, X.; Yang, Y.; Wang, X. Lightweight Image Super-Resolution with Information Multi-Distillation Network. In Proceedings of the 27th ACM International Conference on Multimedia (ACM MM), 2019, pp. 2024–2032. arXiv:1909.11856, <https://doi.org/10.1145/3343031.3351084>.
18. Liu, J.; Tang, J.; Wu, G. Residual Feature Distillation Network for Lightweight Image Super-Resolution. In Proceedings of the European Conference on Computer Vision Workshops (ECCVW), 2020, pp. 41–55. arXiv:2009.11551; AIM 2020 Efficient SR Challenge, [https://doi.org/10.1007/978-3-030-67070-2\\_2](https://doi.org/10.1007/978-3-030-67070-2_2).
19. Kong, F.; Li, M.; Liu, S.; Liu, D.; He, J.; Bai, Y.; Chen, F.; Fu, L. Residual Local Feature Network for Efficient Super-Resolution. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), 2022, pp. 765–775. arXiv:2205.07514; NTIRE 2022 Efficient SR Challenge winner, <https://doi.org/10.1109/CVPRW56347.2022.00092>.
20. Sun, L.; Pan, J.; Tang, J. ShuffleMixer: An Efficient ConvNet for Image Super-Resolution. In Proceedings of the Advances in Neural Information Processing Systems (NeurIPS), 2022, Vol. 35, pp. 17314–17326. arXiv:2205.15175, <https://doi.org/10.52202/068431-1259>.
21. Guo, Y.; Tan, M.; Deng, Z.; Wang, J.; Chen, Q.; Cao, J.; Xu, Y.; Chen, J. Towards Lightweight Super-Resolution with Dual Regression Learning. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **2024**, *46*, 8365–8379. arXiv:2207.07929, <https://doi.org/10.1109/TPAMI.2024.3406556>.
22. Ren, B.; Li, Y.; Mehta, N.; Timofte, R.; et al. The Ninth NTIRE 2024 Efficient Super-Resolution Challenge Report. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)* **2024**, pp. 6595–6631. arXiv:2404.10343, <https://doi.org/10.1109/CVPRW63382.2024.00656>.
23. Wang, X.; Girshick, R.; Gupta, A.; He, K. Non-Local Neural Networks. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2018, pp. 7794–7803. arXiv:1711.07971, <https://doi.org/10.1109/CVPR.2018.00813>.
24. Chen, H.; Wang, Y.; Guo, T.; Xu, C.; Deng, Y.; Liu, Z.; Ma, S.; Xu, C.; Xu, C.; Gao, W. Pre-Trained Image Processing Transformer. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2021, pp. 12294–12305. arXiv:2012.00364, <https://doi.org/10.1109/CVPR46437.2021.01212>.
25. Liu, Z.; Lin, Y.; Cao, Y.; Hu, H.; Wei, Y.; Zhang, Z.; Lin, S.; Guo, B. Swin Transformer: Hierarchical Vision Transformer Using Shifted Windows. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), 2021, pp. 9992–10002. arXiv:2103.14030, <https://doi.org/10.1109/ICCV48922.2021.00986>.
26. Liang, J.; Cao, J.; Sun, G.; Zhang, K.; Van Gool, L.; Timofte, R. SwinIR: Image Restoration Using Swin Transformer. In Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops (ICCVW), 2021, pp. 1833–1844. arXiv:2108.10257, <https://doi.org/10.1109/ICCVW54120.2021.00210>.
27. Zhang, X.; Zeng, H.; Guo, S.; Zhang, L. Efficient Long-Range Attention Network for Image Super-Resolution. In Proceedings of the European Conference on Computer Vision (ECCV), 2022, pp. 649–667. arXiv:2203.06697, [https://doi.org/10.1007/978-3-031-19790-1\\_39](https://doi.org/10.1007/978-3-031-19790-1_39).

28. Zhou, Y.; Li, Z.; Guo, C.L.; Bai, S.; Cheng, M.M.; Hou, Q. SRFormer: Permuted Self-Attention for Single Image Super-Resolution. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), 2023, pp. 12734–12745. arXiv:2303.09735, <https://doi.org/10.1109/ICCV51070.2023.01174>.
29. Chen, X.; Wang, X.; Zhou, J.; Qiao, Y.; Dong, C. Activating More Pixels in Image Super-Resolution Transformer. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2023, pp. 22367–22377. HAT; arXiv:2205.04437, <https://doi.org/10.1109/CVPR52729.2023.02142>.
30. Zheng, L.; Zhu, J.; Shi, J.; Weng, S. Efficient Mixed Transformer for Single Image Super-Resolution. *Engineering Applications of Artificial Intelligence* **2024**, *133*, 108035. arXiv:2305.11403, <https://doi.org/10.1016/j.engappai.2024.108035>.
31. Wang, H.; Chen, X.; Ni, B.; Liu, Y.; Liu, J. Omni Aggregation Networks for Lightweight Image Super-Resolution. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2023, pp. 22378–22387. Omni-SR; arXiv:2304.10244, <https://doi.org/10.1109/CVPR52729.2023.02143>.
32. Choi, H.; Lee, J.; Yang, J. N-Gram in Swin Transformers for Efficient Lightweight Image Super-Resolution. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2023, pp. 2071–2081. NGswin, <https://doi.org/10.1109/CVPR52729.2023.00206>.
33. Liu, J.; Chen, C.; Tang, J.; Wu, G. From Coarse to Fine: Hierarchical Pixel Integration for Lightweight Image Super-Resolution. *Proceedings of the AAAI Conference on Artificial Intelligence* **2023**, *37*, 1666–1674. arXiv:2211.16776, <https://doi.org/10.1609/aaai.v37i2.25254>.
34. Zhang, L.; Li, Y.; Zhou, X.; Zhao, X.; Gu, S. Transcending the Limit of Local Window: Advanced Super-Resolution Transformer with Adaptive Token Dictionary. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2024, pp. 2856–2865. ATD; arXiv:2401.08209, <https://doi.org/10.1109/CVPR52733.2024.00276>.
35. Mei, Y.; Fan, Y.; Zhou, Y.; Huang, L.; Huang, T.S.; Shi, H. Image Super-Resolution with Cross-Scale Non-Local Attention and Exhaustive Self-Exemplars Mining. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2020, pp. 5689–5698. arXiv:2006.01424, <https://doi.org/10.1109/CVPR42600.2020.00573>.
36. Yang, F.; Yang, H.; Fu, J.; Lu, H.; Guo, B. Learning Texture Transformer Network for Image Super-Resolution. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2020, pp. 5790–5799. arXiv:2006.04139, <https://doi.org/10.1109/CVPR42600.2020.00583>.
37. Zhang, A.; Ren, W.; Liu, Y.; Cao, X. Lightweight Image Super-Resolution with Superpixel Token Interaction. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), 2023, pp. 12682–12691. <https://doi.org/10.1109/ICCV51070.2023.01169>.
38. Liu, X.; Liu, J.; Tang, J.; Wu, G. CATANet: Efficient Content-Aware Token Aggregation for Lightweight Image Super-Resolution. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2025, pp. 17902–17912. arXiv:2503.06896, <https://doi.org/10.1109/CVPR52734.2025.01668>.
39. Agustsson, E.; Timofte, R. NTIRE 2017 Challenge on Single Image Super-Resolution: Dataset and Study. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), 2017, pp. 1122–1131. <https://doi.org/10.1109/CVPRW.2017.150>.
40. Bevilacqua, M.; Roumy, A.; Guillemot, C.; Alberi-Morel, M.L. Low-Complexity Single-Image Super-Resolution Based on Nonnegative Neighbor Embedding. In Proceedings of the British Machine Vision Conference (BMVC), 2012, pp. 135.1–135.10. <https://doi.org/10.5244/C.26.135>.
41. Zeyde, R.; Elad, M.; Protter, M. On Single Image Scale-Up Using Sparse-Representations. In Proceedings of the International Conference on Curves and Surfaces, 2012, pp. 711–730. [https://doi.org/10.1007/978-3-642-27413-8\\_47](https://doi.org/10.1007/978-3-642-27413-8_47).
42. Martin, D.; Fowlkes, C.; Tal, D.; Malik, J. A Database of Human Segmented Natural Images and Its Application to Evaluating Segmentation Algorithms and Measuring Ecological Statistics. In Proceedings of the IEEE International Conference on Computer Vision (ICCV), 2001, Vol. 2, pp. 416–423. <https://doi.org/10.1109/ICCV.2001.937655>.



43. Huang, J.B.; Singh, A.; Ahuja, N. Single Image Super-Resolution from Transformed Self-Exemplars. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2015, pp. 5197–5206. <https://doi.org/10.1109/CVPR.2015.7299156>.  
734  
735  
736
44. Matsui, Y.; Ito, K.; Aramaki, Y.; Fujimoto, A.; Ogawa, T.; Yamasaki, T.; Aizawa, K. Sketch-Based Manga Retrieval Using Manga109 Dataset. *Multimedia Tools and Applications* **2017**, *76*, 21811–21838. <https://doi.org/10.1007/s11042-016-4020-z>.  
737  
738  
739
45. Wang, Z.; Bovik, A.C.; Sheikh, H.R.; Simoncelli, E.P. Image Quality Assessment: From Error Visibility to Structural Similarity. *IEEE Transactions on Image Processing* **2004**, *13*, 600–612. <https://doi.org/10.1109/TIP.2003.819861>.  
740  
741  
742

**Disclaimer/Publisher’s Note:** The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.  
743  
744  
745  
746